

Network Science Analytics

Option Applied Math and M.Sc. in DSBA

Lecture 5A

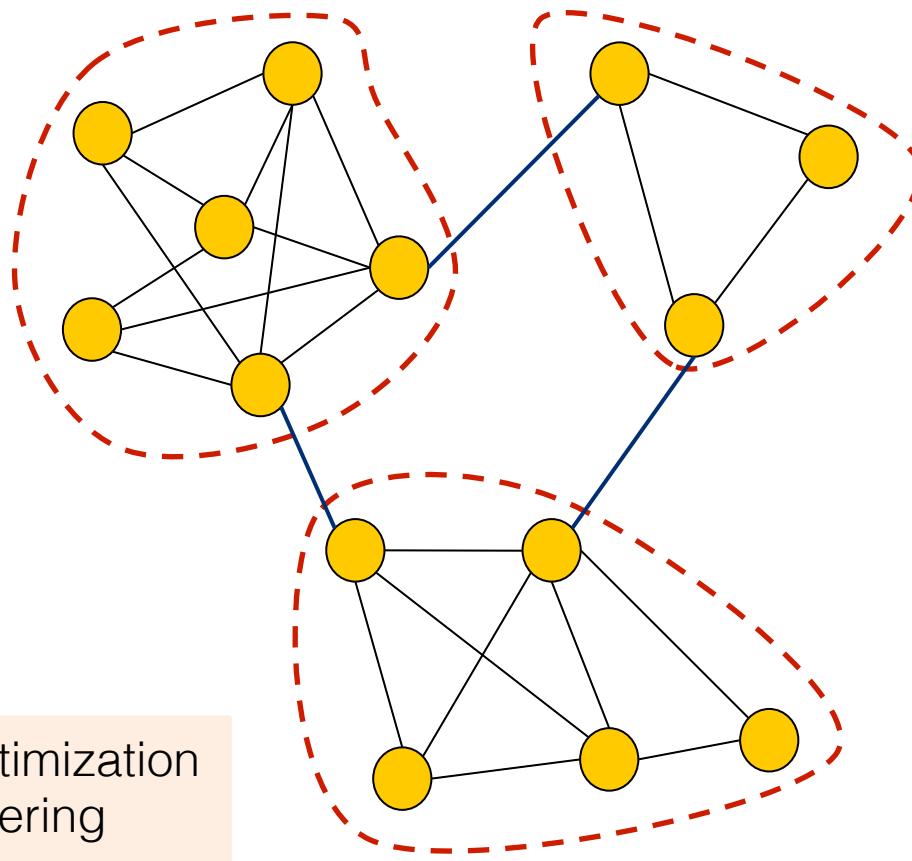
Link prediction in networks

Fragkiskos Malliaros

Friday, March 2, 2018

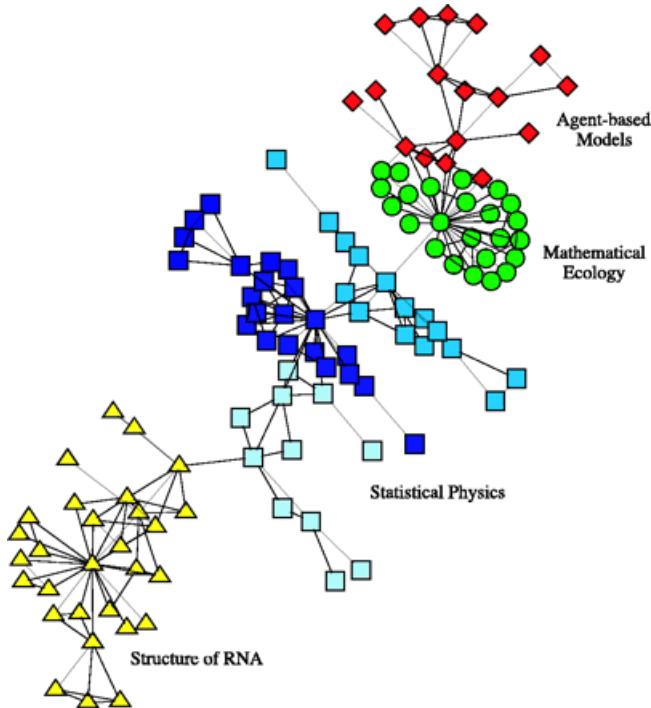
Community structure of large-scale graphs

Community Structure

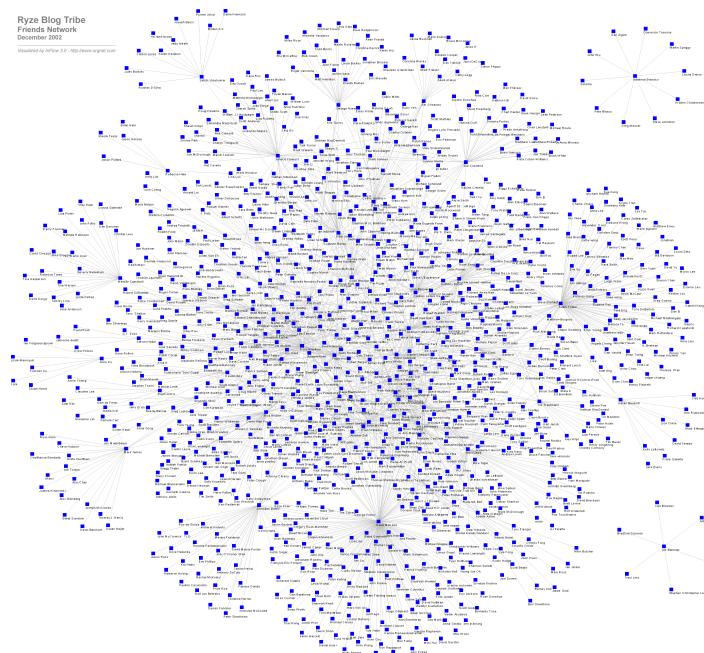


- Modularity optimization
- Spectral clustering
- ...

Community Structure in Small vs. Large Graphs



Small scale collaboration
network (Newman)



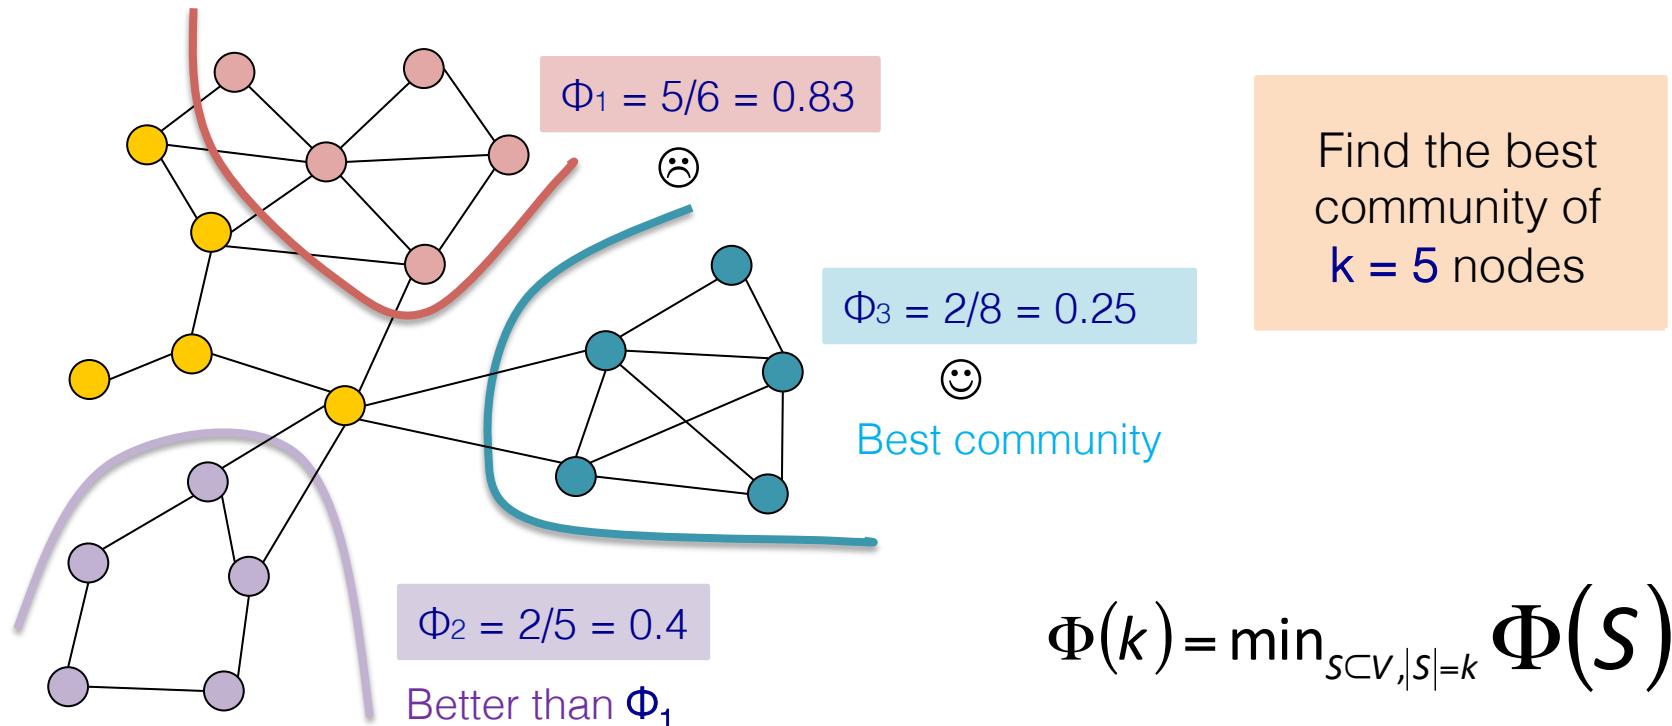
Blog network
<http://www.ryze.com>

Examine the Structural Differences

- Use **conductance $\Phi(S)$** as a community evaluation measure
 - Smaller value for conductance implies better community-like properties [Leskovec et al. '09]

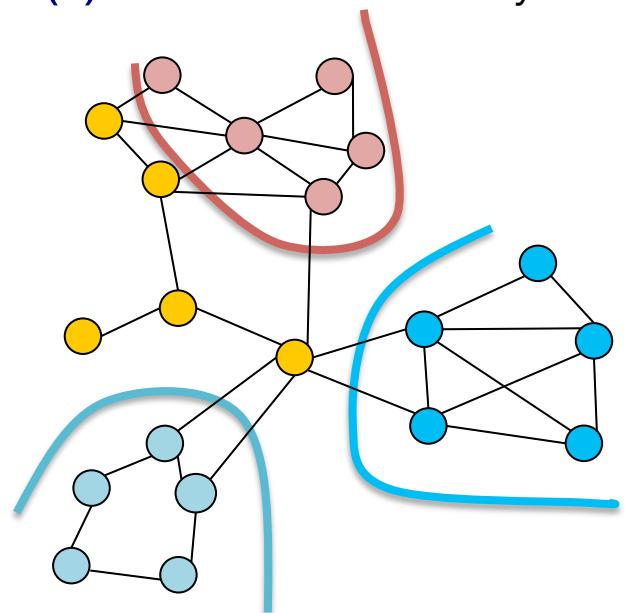
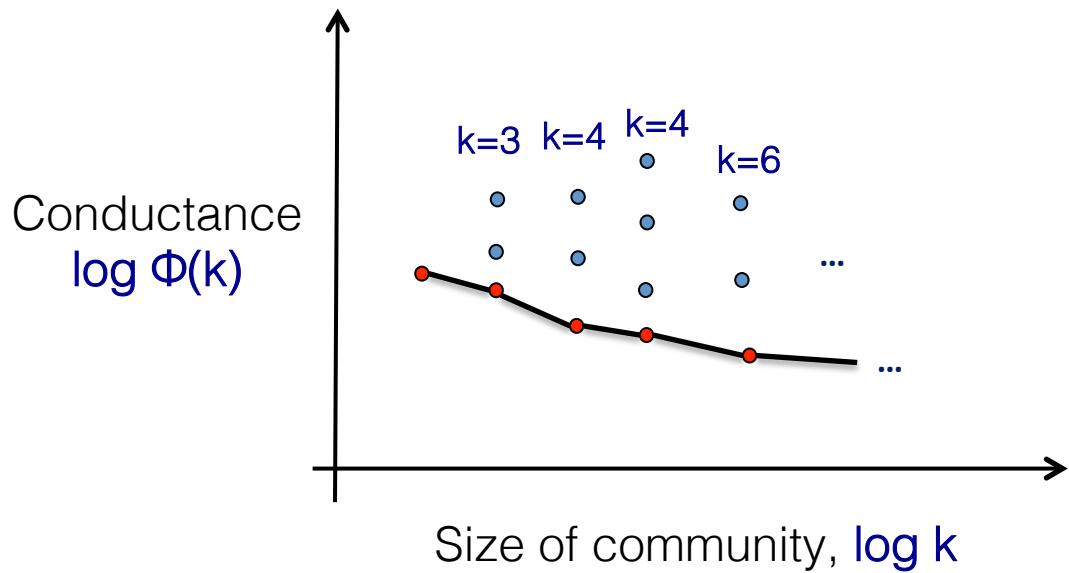
$$\Phi(S) = \# \text{ outgoing edges} / \# \text{ edges within}$$

lower is better



Network Community Profile plot

- Network Community Profile (NCP) plot [Leskovec et al. '09]
 - Plot the best conductance score (minimum) $\Phi(k)$ for each community size k

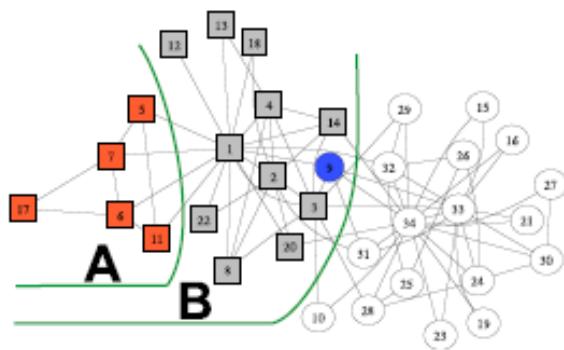


NCP plot of real graphs

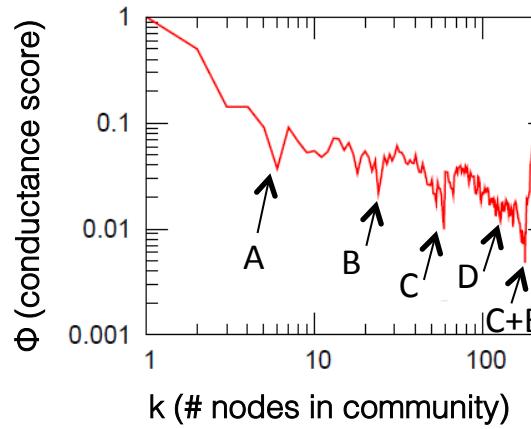
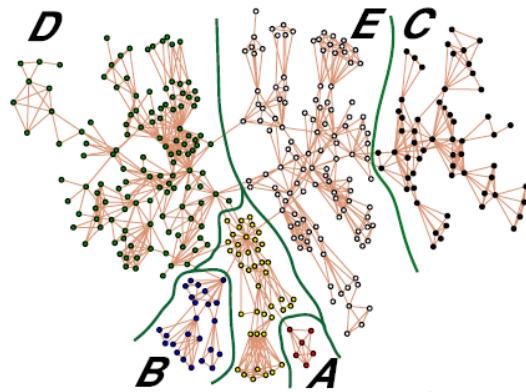
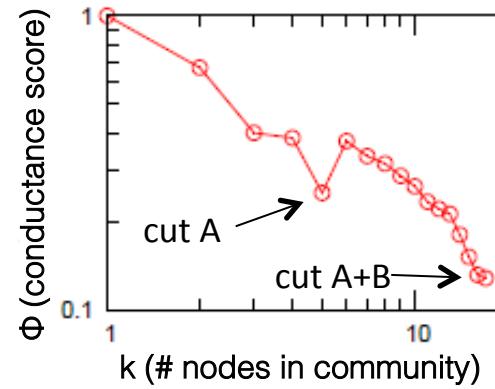
?

NCP Plot Examples

Small scale networks

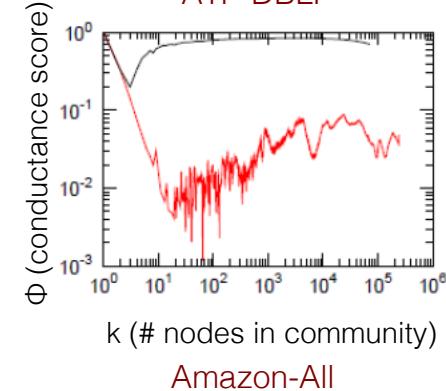
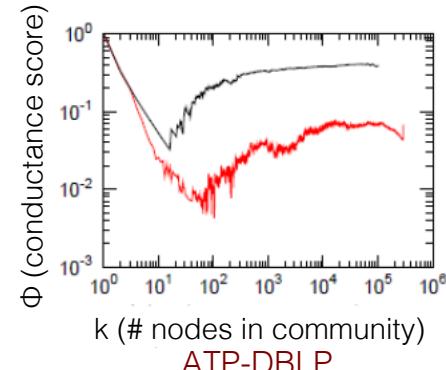
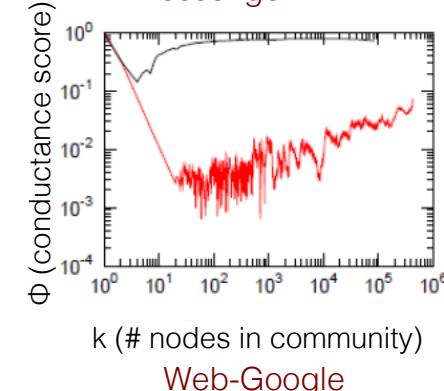
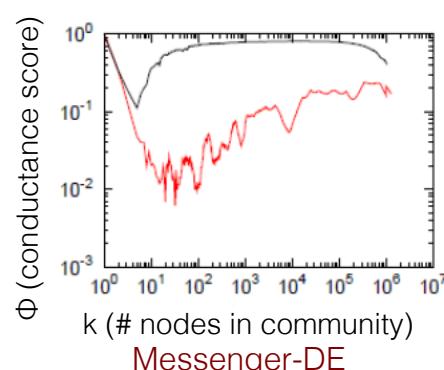
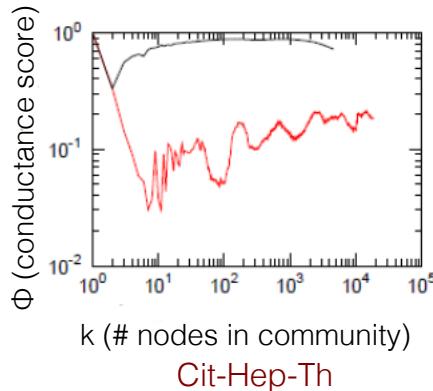
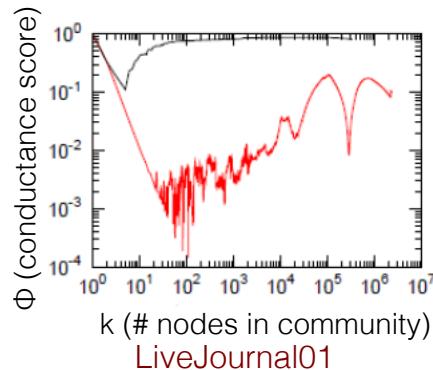


Zachary's karate club social network



Newman's collaboration network

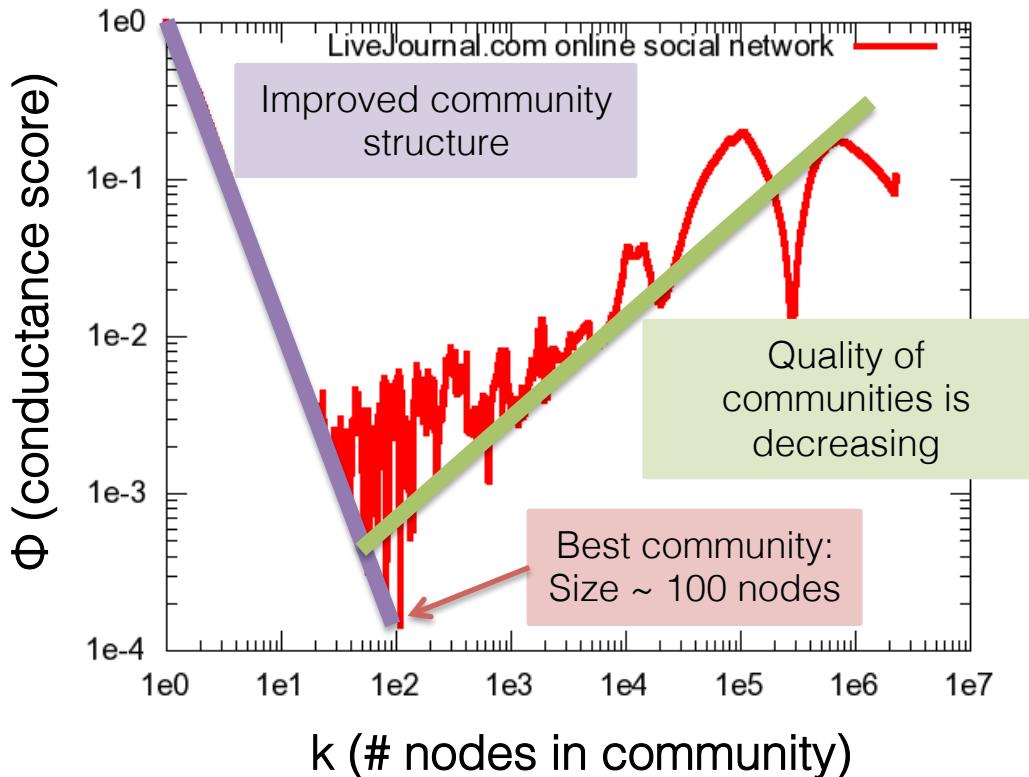
NCP Plot of Large Real Graphs



Any common property?

Large scale
networks

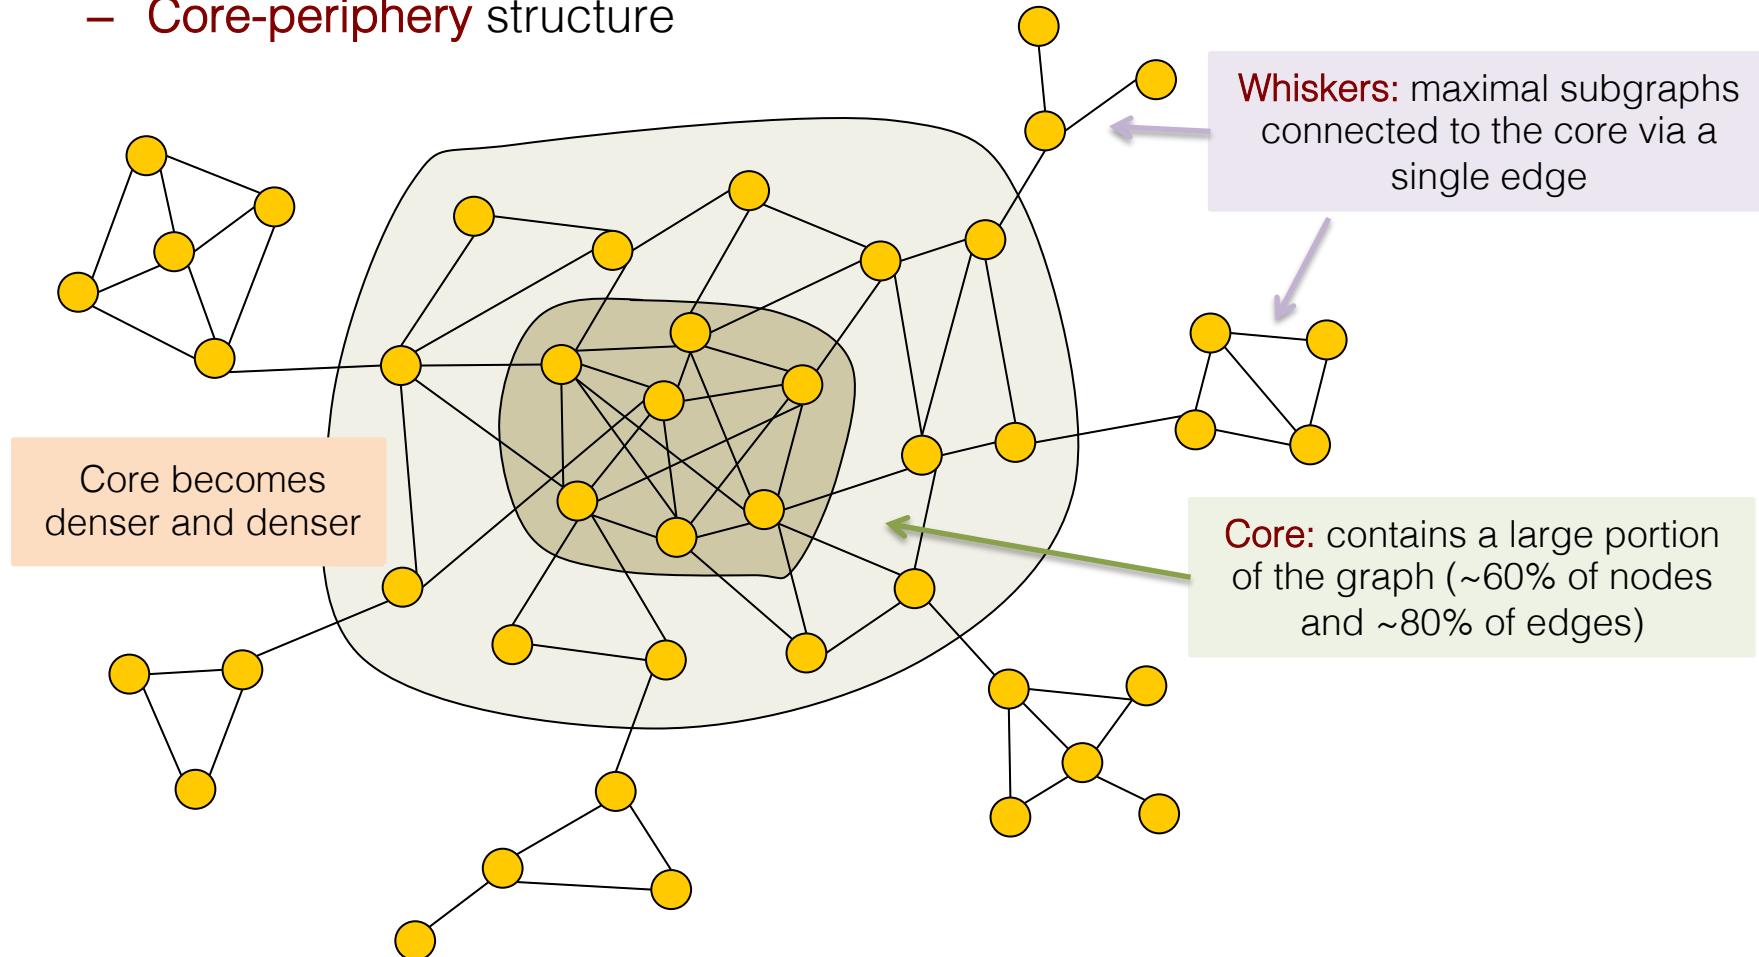
Observation in Large Graphs



LiveJournal social network
 $|V| = 5M, |E| = 42M$

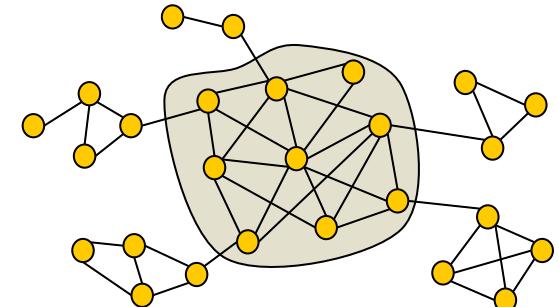
Explanation: Core-Periphery Structure

- How can we explain the observed structure of large graphs?
 - **Core-periphery** structure



Core-Periphery Structure

- Core-periphery structure
 - Core
 - Whiskers
- Whiskers
 - Non-trivial structure → more than random (shape and size)
- Q: What is happening if we remove whiskers (periphery) from graphs?
 - Almost nothing. The whiskers are replaced by 2-whiskers (subgraphs connected to the core with 2 edges)
- The core itself has core-periphery structure
- Important point: Whiskers are also responsible for the best communities in large graphs (lowest point of NCP plot)



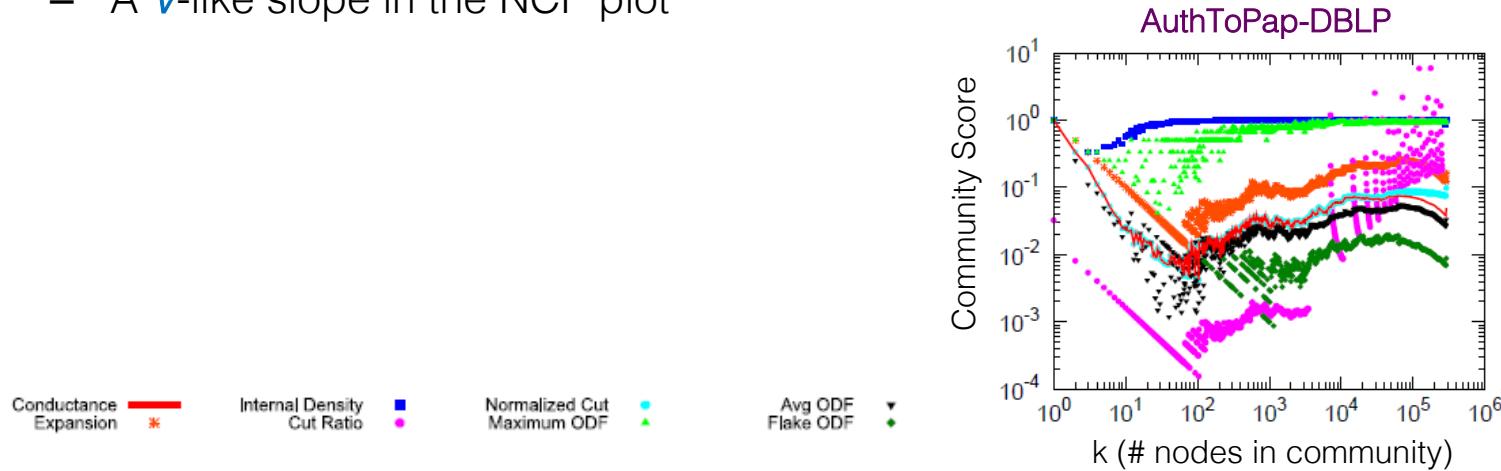
Match the observations made using the properties of the Kronecker graphs

Similar Structural Observations

- **Jellyfish model** for the Internet topology [Tauro et al. '01]
- **Min-cut** plots [Chakrabarty et al. '04]
 - Perform min-cut recursively
 - Plot the relative size of the minimum cut
- **Robustness** of large scale social networks [Malliaros, Megalooikonomou, Faloutsos '12, '15]
 - Robustness estimation based on the expansion properties of graphs
 - Social networks are expected to have **low robustness** due to the existence of communities → the (small number of) inter-community edges will act as bottlenecks
 - Large scale social graphs tend to be extremely robust
 - **Structural differences** (in terms of robustness and community structure) between **different scale graphs**

Clustering Algorithms and Objective Criteria

- **Question 1:** Is the observed property an effect of the used community detection algorithm (Metis + flow based method)?
 - A: No. The qualitative shape of the NCP plot is the same, regardless of the community detection algorithm [Leskovec et al. '09]
- **Question 2:** Is the observed property an effect of the **conductance** community evaluation measure?
 - A: No. All the objective criteria that based on both internal and external connectivity, show an almost similar qualitatively behavior [Leskovec et al. '10]
 - A V-like slope in the NCP plot



Conclusions

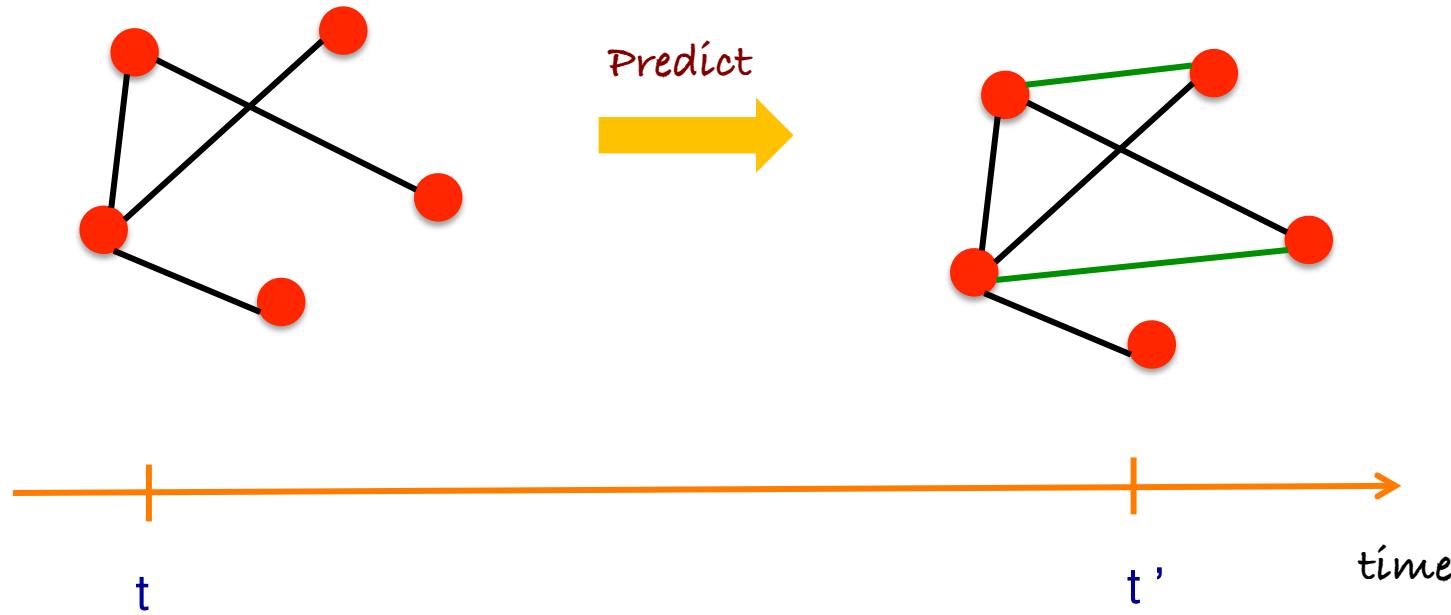
- Large scale real-world graphs
 - Core-periphery structure
 - No large, well defined communities
 - Structural differences between different scale graphs
- Community detection algorithms should take into account these structural observations
 - Whiskers correspond to the best (conductance-based) communities
 - Need larger high-quality clusters?
 - **Bag of whiskers:** union of disjoint (disconnected) whiskers are mainly responsible for the best high-quality clusters of larger size (above 100)

Survey Articles

- S. Fortunato. **Community detection in graphs.** Physics Reports 486, 75-174, 2010
- S. E. Schaeffer. **Graph clustering.** Computer Science Review, 1(1): 27–64, 2007
- F. D. Malliaros and M. Vazirgiannis. **Clustering and community detection in directed networks: A survey.** Physics Reports 533, 95-142, 2013
- S. Fortunato and D. Hric. **Community detection in networks: a user guide.** Physics Reports 659, 1-44, 2016
- S. Papadopoulos, Y. Kompatsiaris, A. Vakali, and P. Spyridonos. **Community detection in Social Media: Performance and application considerations.** Data Min Knowl Disc, 24:515–554, 2012

Link prediction

Link Prediction



In this lecture

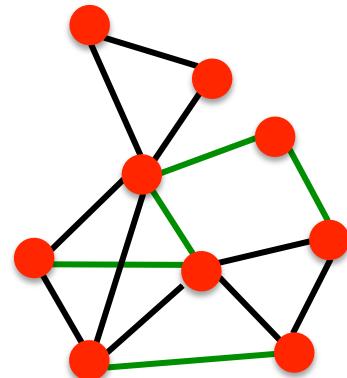
- Unsupervised link prediction
- Supervised link prediction
- Learning to predict links based on supervised random walks

Motivation

- Recommending **new friends** in online social networks
- Suggesting **interactions** between the members of a company/organization that are external to the hierarchical structure of the organization itself
- Suggesting **collaborations** between researchers based on co-authorship
- Predicting **connections** between members of terrorist organizations who have not been directly observed to work together
- Overcoming the data-sparsity problem in **recommender systems** using collaborative filtering

Link Prediction in Networks (1/2)

- The link prediction task
 - Given $G[t_0, t_0']$ a graph on edges up to time t_0' , output a ranked list L of links (not in $G[t_0, t_0']$) that are predicted to appear in $G[t_1, t_1']$
- Evaluation
 - $n=|E_{\text{new}}|$: # of new edges that appear during the test period $[t_1, t_1']$
 - Take top n elements of list L and count the correct edges



$G[t_0, t_0']$
 $G[t_1, t_1']$

Link Prediction in Networks (2/2)

- Prediction for a subset of nodes
- Two parameters: k_{training} and k_{test}
- **Core:** all nodes that are incident to at least k_{training} edges in $G[t_0, t'_0]$, and at least k_{test} edges in $G[t_1, t'_1]$
 - Networks evolve over time
 - Predict only edges whose endpoints appear in both the **training** and **test** intervals
- Predict new edges between the nodes in **Core**

Problem Formulation (2/2)

	training period			Core		
	authors	papers	collaborations ¹	authors	$ E_{old} $	$ E_{new} $
astro-ph	5343	5816	41852	1561	6178	5751
cond-mat	5469	6700	19881	1253	1899	1150
gr-qc	2122	3287	5724	486	519	400
hep-ph	5414	10254	47806	1790	6654	3294
hep-th	5241	9498	15842	1438	2311	1576

Example: Predict links in an evolving collaboration network

- $t_0 = 1994, t'_0 = 1996$: training interval $\rightarrow [1994, 1996]$
- $t_1 = 1997, t'_1 = 1999$: test interval $\rightarrow [1997, 1999]$

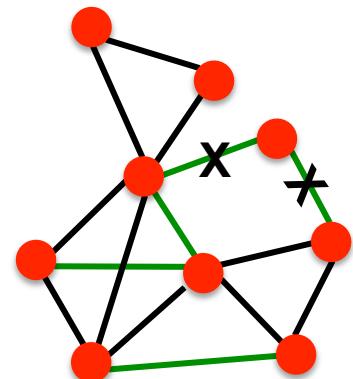
- $G_{collab} = (V, E_{old}) = G[1994, 1996]$
- E_{new} : authors in V that co-author a paper during the test interval but not during the training interval

$k_{\text{training}} = 3, k_{\text{test}} = 3$: Core consists of all authors who have written at least 3 papers during the training period and at least 3 papers during the test period

Goal: predict E_{new}

Link Prediction - Evaluation

- For each pair of nodes (x, y) , compute score $c(x, y)$
 - For example, $c(x, y)$ could be the # of common neighbors of x and y
- Sort pairs (x, y) by the decreasing score $c(x, y)$
 - Only consider/predict edges where both endpoints are in the core
- Predict the top n pairs of new links
- See which of those links actually appear in $G[t_1, t'_1]$



Methods of Link Prediction

How to assign score $c(x, y)$ between two nodes x and y ?

Some form of **similarity** or **node proximity** of x and y

- Many different ways to formalize this
- Use only graph features

Methods

- Neighborhood-based (# of shared neighbors)
- Network proximity-based methods (paths between x and y)

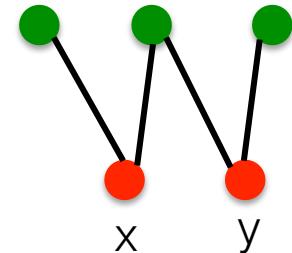
Neighborhood-based Methods (1/4)

- Let $\Gamma(x)$ be the set of neighbors of x in G_{old}
- Methods
 - Common neighbors overlap
 - Jaccard coefficient
 - Adamic/Adar index
 - Preferential attachment

Intuition: The larger the overlap of the neighbors of two nodes, the more likely the nodes to be linked in the future

Neighborhood-based Methods (2/4)

Let $\Gamma(x)$ be the set of neighbors of x in G_{old}



Common neighbors:

$$c(x, y) = |\Gamma(x) \cap \Gamma(y)|$$

A : adjacency matrix
 $A_{x,y}^2$: Number of different paths of length 2

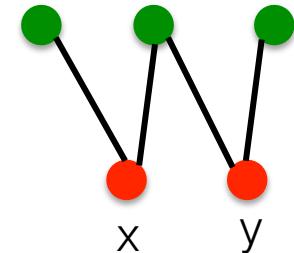
Jaccard coefficient:

$$c(x, y) = \frac{|\Gamma(x) \cap \Gamma(y)|}{|\Gamma(x) \cup \Gamma(y)|}$$

The probability that both x and y have common neighbors

Neighborhood-based Methods (3/4)

Let $\Gamma(x)$ be the set of neighbors of x in G_{old}



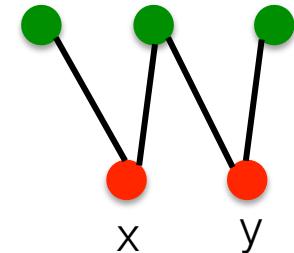
Adamic/Adar:

$$c(x, y) = \sum_{z \in \Gamma(x) \cap \Gamma(y)} \frac{1}{\log |\Gamma(z)|}$$

Assigns large weights to common neighbors z of x and y which themselves have few neighbors (weights rare features more heavily)

Neighborhood-based Methods (4/4)

Let $\Gamma(x)$ be the set of neighbors of x in G_{old}



Preferential attachment:

$$c(x, y) = |\Gamma(x)| \cdot |\Gamma(y)|$$

of neighbors = degree

Based on the intuition that the probability that a new edge has node x as its endpoint is proportional to $|\Gamma(x)|$, i.e., nodes prefer to form ties with ‘popular’ nodes

Proximity-based Methods (1/4)

Intuition: The “closer” two nodes are in the network, the more likely is to be linked in the future

- Methods
 - Based on **shortest path length** between **x** and **y**
 - Based on **all paths** between **x** and **y**
 - Katz _{β} measure
 - Random walk-based
 - Hitting time
 - Commute time
 - Rooted PageRank
 - SimRank

Shortest Path-Based

For nodes x, y

$c(x, y) = (\text{negated}) \text{ length of shortest path}$
between x and y

*Some further
normalization may
needed*

If there are more than n pairs of nodes at the shortest path length ℓ ,
order them at random

Ensemble of All Paths (1/2)

Katz _{β} measure:

$$c(x, y) = \sum_{\ell=1}^{\infty} \beta^\ell \cdot |\text{paths}_{x,y}^{(\ell)}|$$

Set of all paths of length ℓ from x to y

- $0 < \beta < 1$ is a parameter of the predictor, exponentially damped to count short paths more heavily
- Small β yields predictions much like common neighbors

Ensemble of All Paths (2/2)

Katz _{β} measure:

$$c(x, y) = \sum_{\ell=1}^{\infty} \beta^{\ell} \cdot |\text{paths}_{x,y}^{(\ell)}|$$

Set of all paths of length ℓ from x to y

$$\sum_{\ell=1}^{\infty} \beta^{\ell} \cdot |\text{paths}_{x,y}^{(\ell)}| = \beta A_{xy} + \beta^2 (A^2)_{xy} + \beta^3 (A^3)_{xy} + \dots$$

Closed form: $(I - \beta A)^{-1} - I$ *Matrix of scores*

- Unweighted version: $\text{path}_{x,y}^{(1)} = 1$ (edge with ‘weight = 1’)
- Weighted version: $\text{path}_{x,y}^{(1)} = \# \text{ times } x \text{ and } y \text{ have collaborated}$

Random Walk-based Methods (1/2)

Consider a random walk on G_{old} that starts at x and iteratively moves to a neighbor of x chosen uniformly at random from $\Gamma(x)$

- **Hitting time** $H_{x,y}$ (from x to y): the expected number of steps it takes for the random walk starting at x to reach y

$$c(x, y) = -H_{x,y} \quad \begin{matrix} \text{Not symmetric,} \\ \text{in general} \end{matrix}$$

- **Commute time** $\text{Comm}(x, y)$ (from x to y): the expected number of steps to travel from x to y and from y to x

$$c(x, y) = - (H_{x,y} + H_{y,x})$$

Random Walk-based Methods (2/2)

- The hitting time and commute time measures are sensitive to parts of the graph far away from x and y
 - Periodically reset the walk
- Random walk on G_{old} that starts at x and has a probability a of returning to x at each step
- Rooted (Personalized) PageRank
 - Starts from x
 - With probability $(1 - a)$ moves to a random neighbor
 - With probability a returns to x

$c(x, y) = \text{stationary probability of } y \text{ in a rooted PageRank}$

SimRank

- **Intuition:** Two objects are similar, if they are related to similar objects
- Two objects x and y are similar, if they are related to objects a and b respectively and a and b are themselves similar

$$\text{sim}(x, y) = \frac{\gamma}{|\Gamma(x)| \cdot |\Gamma(y)|} \sum_{a \in \Gamma(x)} \sum_{b \in \Gamma(y)} \text{sim}(a, b)$$

Recursive definition

computation

$$\text{sim}_0(a, b) = \begin{cases} 1, & \text{if } a = b \\ 0, & \text{if } a \neq b \end{cases}$$

Repeat iteration until convergence

Expresses the average similarity between neighbors of x and neighbors of y

$$c(x, y) = \text{sim}(x, y)$$

[Jeh and Widom '02]

Evaluation and results

How to Evaluate the Prediction?

- Each link predictor p outputs a ranked list L_p of pairs in $V \times V - E_{\text{old}}$: predicted new edges in decreasing order of confidence
- Here we focus on **Core** of the graph (degree $k > 3$), thus

$$E_{\text{new}}^* = E_{\text{new}} \cap (\text{Core} \times \text{Core})$$

- Evaluation method: size of intersection of
 - The first n edge predictions from L_p that are in $\text{Core} \times \text{Core}$ and
 - The actual set E_{new}

How many of the (relevant) top- n predictions are correct (precision?)

Evaluation: Baseline Predictor

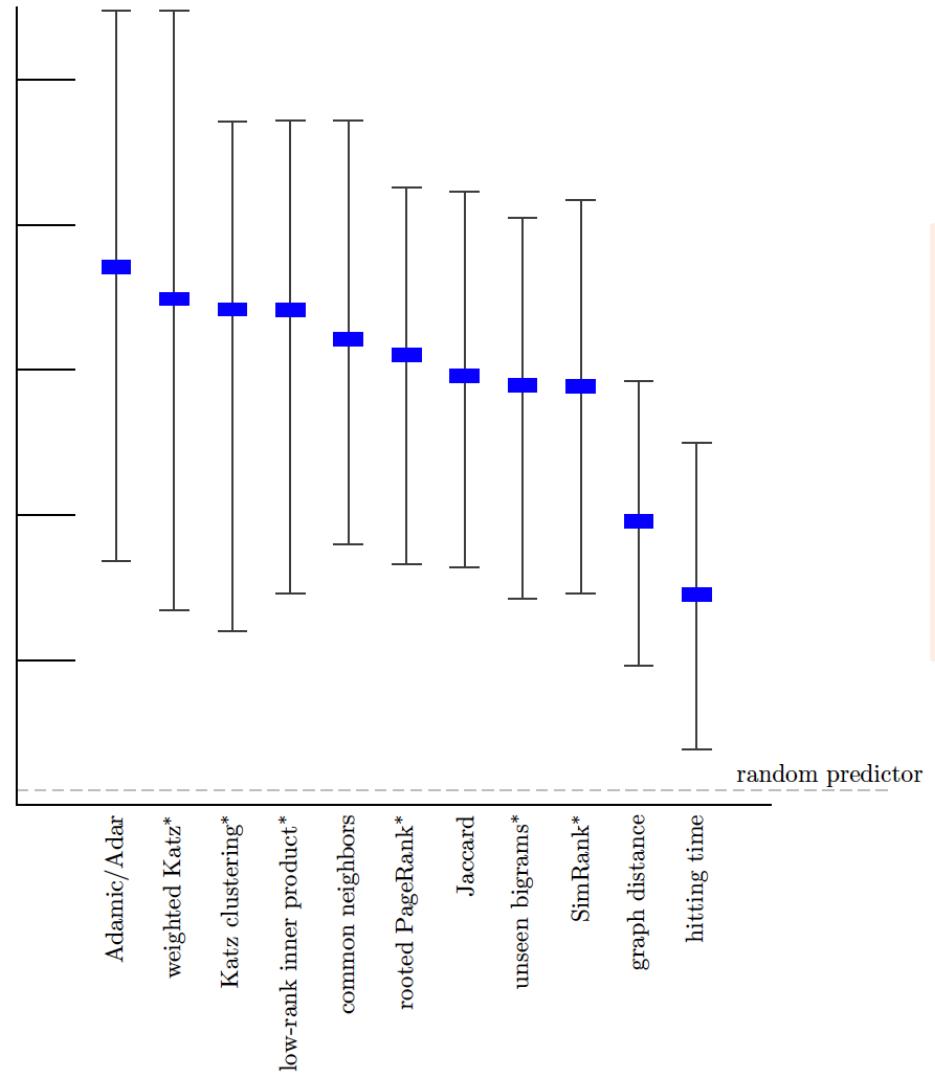
- Prediction accuracy will be measured in terms of relative improvement over a **random predictor**
 - Why? First application was in co-authorship networks
 - Many collaborations form (or fail to form) for reasons outside the scope of the networks
 - The raw accuracy (in the unsupervised case) is very low
- **Random Predictor:** randomly selects pairs of authors who did not collaborate (i.e., there is no edge) in the training interval
 - Probability that a random prediction is correct:

$$\frac{|E_{new}|}{\binom{|\text{Core}|}{2} - |E_{old}|}$$

Random predictions are correct with prob. between 0.15% (cond-mat) to 0.48% (astro-ph)

Relative Average Performance

Relative performance ratio versus random predictions



- Relative avg. performance of various predictors vs. the random predictor
- Avg. ratio over 5 datasets
- The Adamic/Adar and Katz measures perform very well

Discussion and Extensions

- Improve **performance**. Even the best (Katz clustering on gr-qc) is correct on only about 16% of its prediction
- Improve **efficiency** on very large networks (approximation of distances)
- Treat more **recent links** (e.g., collaborations) as more important
- **Additional information** (paper titles, author institutions, etc.) latently present in the graph
- Why not using those measures as features in a model that performs the predictions?

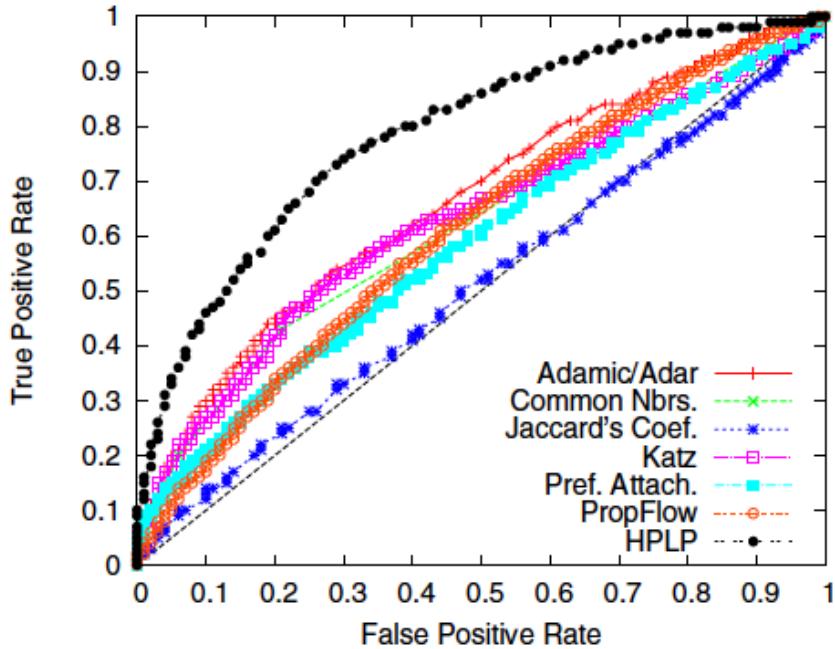
Supervised Link Prediction

- The same “similarity” measures can be used as features for supervised link prediction
 - Training and test datasets
 - Similar to a binary classification task: ‘**class 1**’ (edge exists); ‘**class 0**’ (edge doesn’t exist)
 - Use **ROC curve** (false positive rate vs. true positive rate) to evaluate the predictions
 - Weak point: have to deal with **class imbalance**

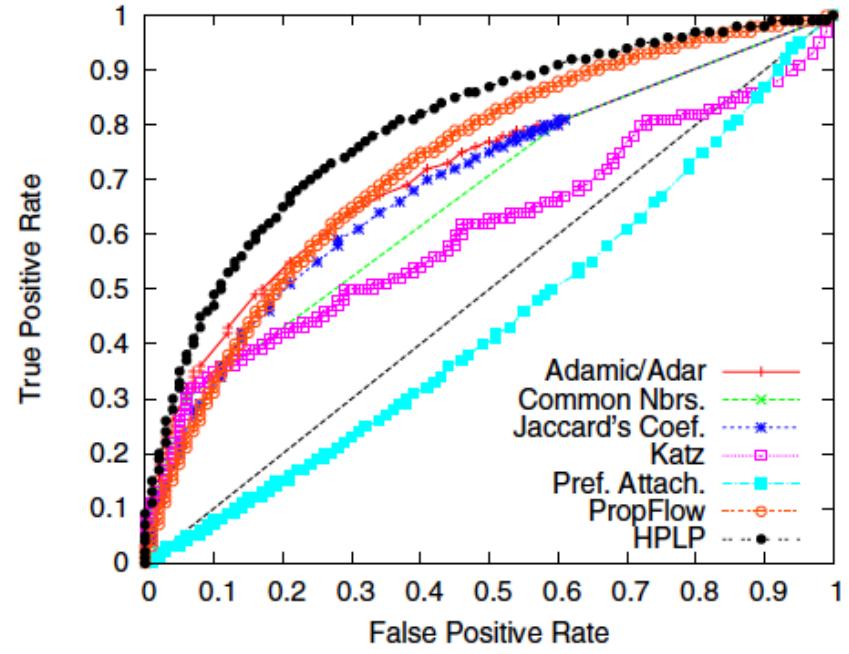
Features

Name	Parameters	HPLP	HPLP+
In-Degree(i)	-	✓	✓
In-Volume(i)	-	✓	✓
In-Degree(j)	-	✓	✓
In-Volume(j)	-	✓	✓
Out-Degree(i)	-	✓	✓
Out-Volume(i)	-	✓	✓
Out-Degree(j)	-	✓	✓
Out-Volume(j)	-	✓	✓
Common Nbrs(i,j)	-	✓	✓
Max. Flow(i,j)	$l = 5$	✓	✓
Shortest Paths(i,j)	$l = 5$	✓	✓
PropFlow(i,j)	$l = 5$	✓	✓
Adamic/Adar(i,j)	-		✓
Jaccard's Coef(i,j)	-		✓
Katz(i,j)	$l = 5, \beta = 0.005$		✓
Pref Attach(i,j)	-		✓

Experimental Results



Cond-Mat



Phone

- Random forests classifier
- HPLP: method that combines multiple features

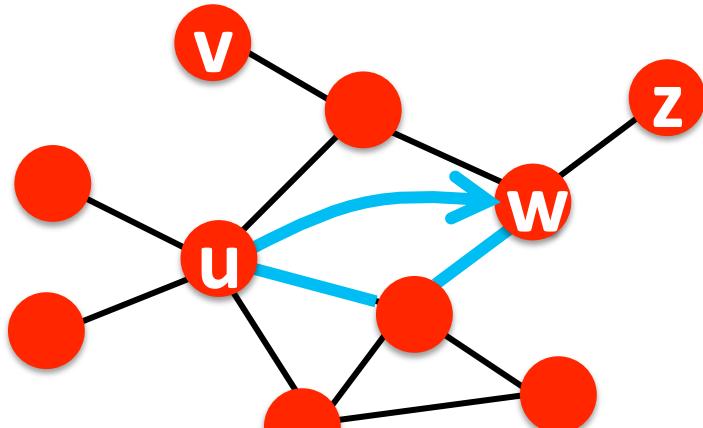
[Lichtenwalter et al. '10]

LPMade: Link Prediction Made Easy
Source: <https://www3.nd.edu/~dial/software/>

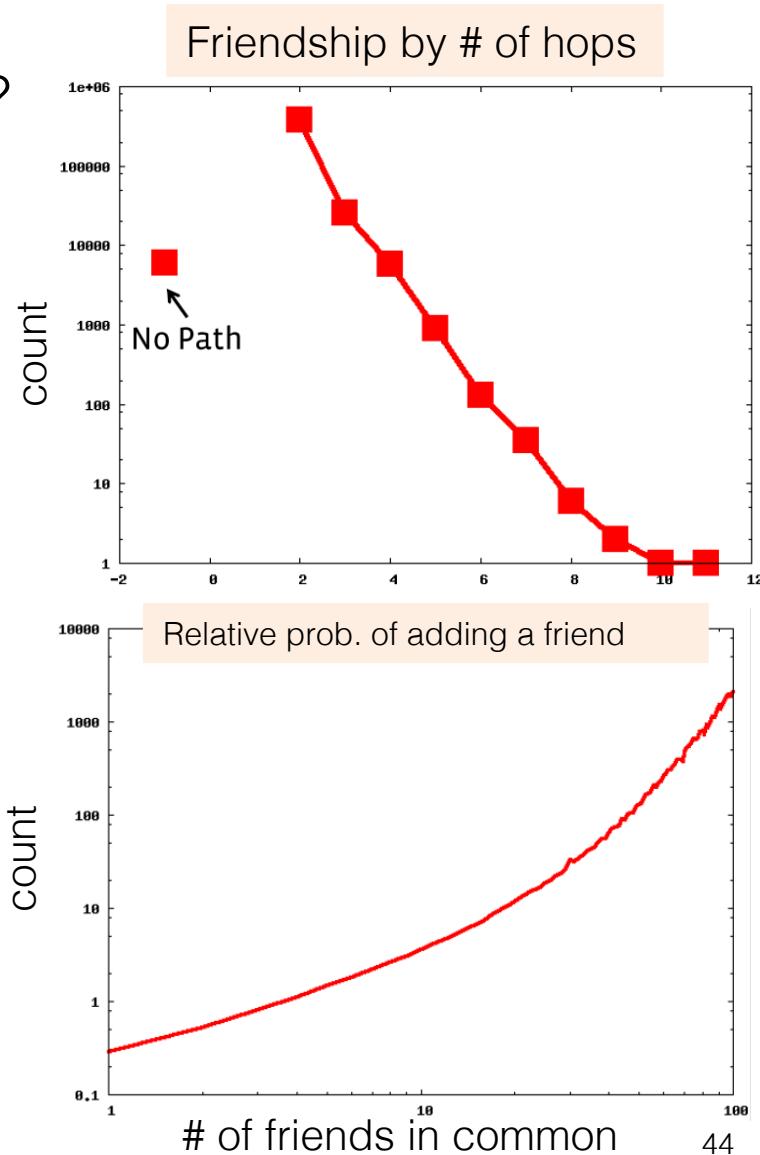
Supervised random walks for link prediction

Supervised Link Prediction - Motivation

- Can we learn to predict new friends?
 - Facebook's People You May Know
 - Let's look at the FB data
 - 92% of new friendships on FB are friend-of-a-friend
 - More mutual friends help



[Backstrom and Leskovec '11]

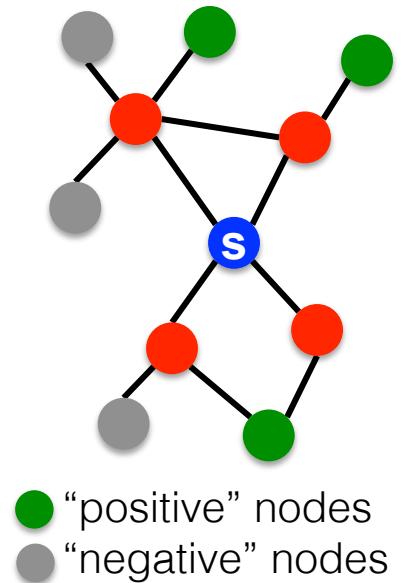


Supervised Link Prediction - Overview

- Proximity measure + classification approach [previous part]
 - Class imbalance
 - Which proximity measures to use
- Rank the nodes: higher score to nodes that **s** will create link to
 - PageRank, Random walk with restarts, etc.
 - Node/edge features are not taken into account
- In this part: use the idea of **Supervised Random Walks**
 - Combine information about the **network structure** and **node/edge attributes** (e.g., age, gender, home town, school – profile information)
 - Use this information to **guide** the random walk
 - Learn a function that assigns **strengths** (weights) to the edges ...
 - i.e., random walk transition probabilities
 - ... such that a random walker is more likely to visit the nodes to which new links will be created in the future

Supervised Link Prediction (1/2)

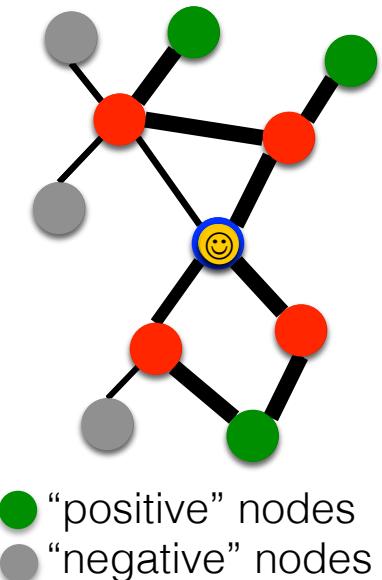
- **Goal:** Recommend a list of possible friends
- Supervised ML setting
 - **Labeled training** examples:
 - For every user s , we have a list of other users she will create links to in the future: $\{d_1, \dots, d_k\}$
 - Use FB data from May 2012 and $\{d_1, \dots, d_k\}$ are the new friendships you created since then
 - These are the **positive** training examples
 - Use all other users as **negative** examples
 - **Task:** For a given node s , score nodes $\{d_1, \dots, d_k\}$ higher than any other nodes in the network



Green nodes are the nodes to which s creates links in the future

Supervised Link Prediction (2/2)

- How to combine node/edge features and the network structure?
 - Estimate **strength** of each friendship (u, v) using:
 - Profile of user u , profile of user v
 - Interaction history of users u and v
 - This creates a **weighted graph**
 - Do **Personalized PageRank** from s and measure the “**proximity**” (the visiting prob.) of any other node w from s
 - Sort nodes w by decreasing **proximity**

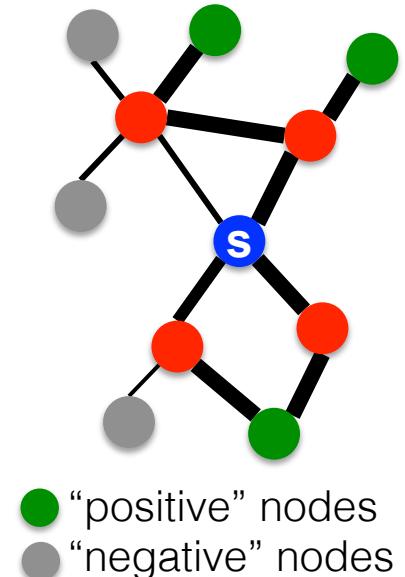


Supervised Random Walks

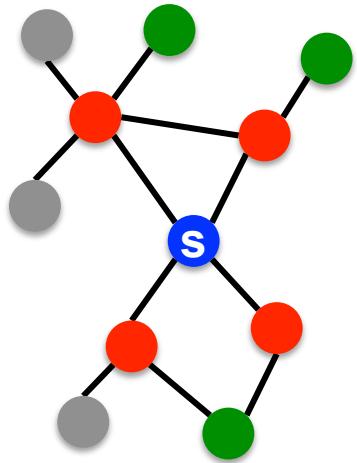
- Let s be the starting node
- Let $f_\beta(u, v)$ be a function that assigns strength a_{uv} to edge (u, v)

$$\alpha_{uv} = f_\beta(u, v) = \exp\left(-\sum_i \beta_i \cdot x_{uv}(i)\right)$$

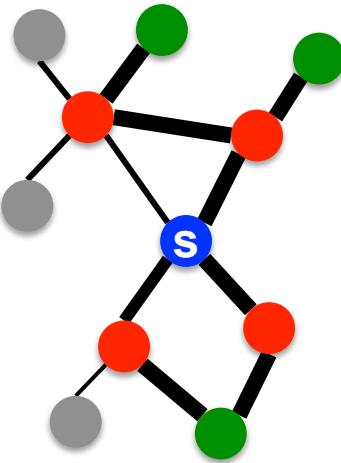
- x_{uv} is feature vector of (u, v)
 - Features of node u
 - Features of node v
 - Features of edge (u, v)
- [Note: β is the weight vector that we will estimate later]
- Do random walk with restarts from s where the transitions are made based on the edge strengths a_{uv}



SRW: Prediction



Network



Set edge
strengths
 $\alpha_{uv} = f_\beta(u, v)$

Random Walk with
Restarts on the
weighted graph.
Each node w has a
PageRank proximity p_w



Sort nodes w by the
decreasing PageRank
score p_w



Recommend top k
nodes with the highest
proximity p_w to node s

- How do we estimate edge strengths?
 - How to estimate parameters β in $f_\beta(u, v)$
- Idea: Set β such that it (correctly) predicts the known future links

Personalized PageRank

- α_{uv} is the strength of edge (u, v)

- Random walk transition matrix

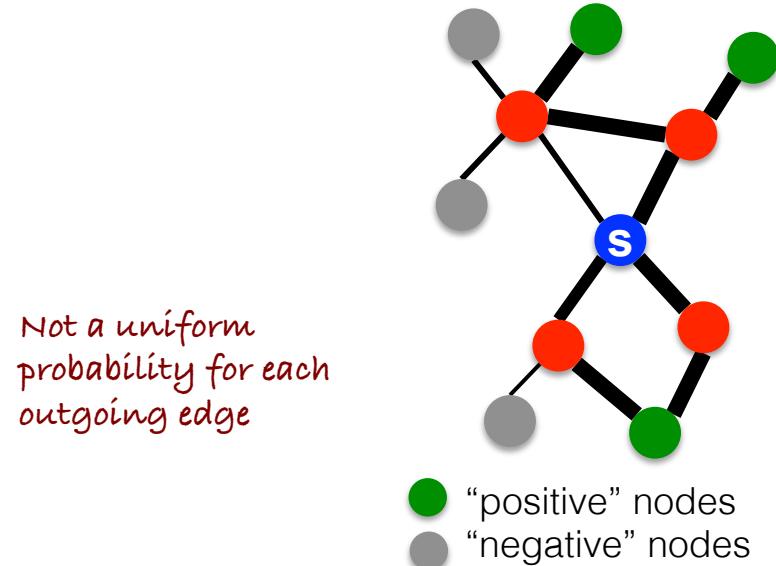
$$Q'_{uv} = \begin{cases} \frac{\alpha_{uv}}{\sum_w \alpha_{uw}}, & \text{if } (u, v) \in E \\ 0, & \text{otherwise} \end{cases}$$

- PageRank transition matrix

$$Q_{ij} = (1 - \gamma)Q'_{ij} + \gamma \mathbf{1}(j = s) \quad \gamma: \text{restart probability of PageRank}$$

- Compute PageRank vector $\mathbf{p} = \mathbf{p}^T Q$

- Rank nodes w by decreasing p_w



The Optimization Problem

- Positive examples $D = \{d_1, \dots, d_k\}$
- Negative examples $L = \{\text{other nodes}\}$
- What do we want?

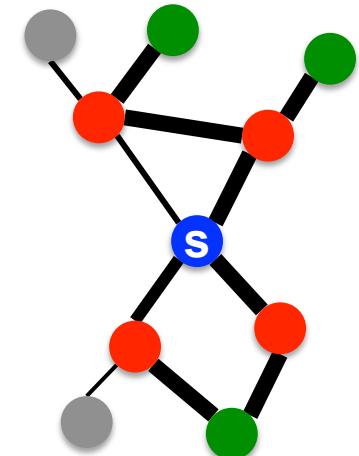
$$\min_{\beta} F(\beta) = \|\beta\|^2$$

Prefer small weights β
to prevent overfitting

such that: $\forall d \in D, l \in L : p_l < p_d$

Every positive example (list D) has to
have higher PageRank score than
every negative example (list L)

- Note:
 - Exact solution to this problem may not exist
 - So we make the constraints “soft” (i.e., optional) - introduce a loss function



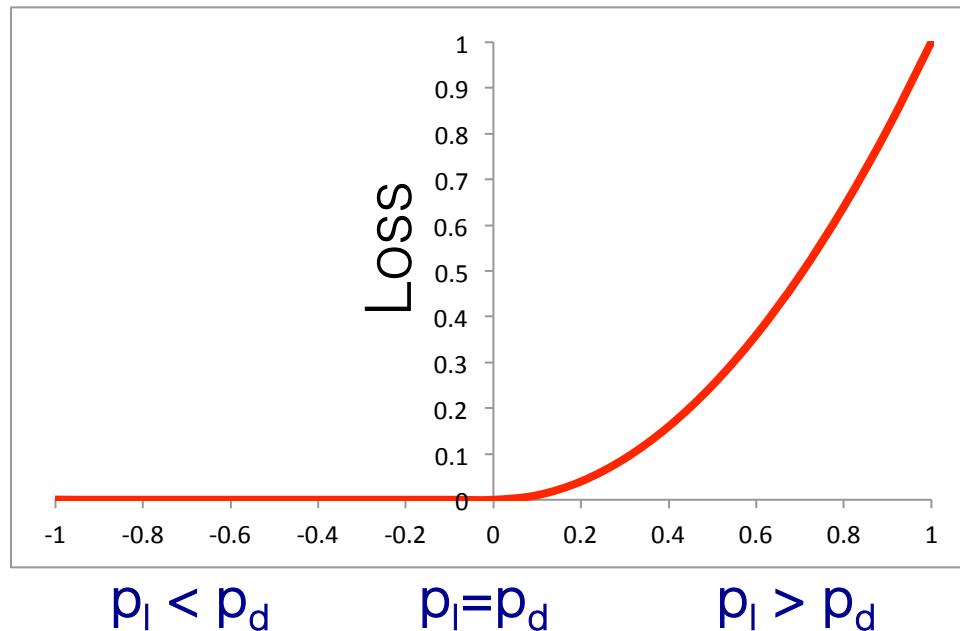
● “positive” nodes
● “negative” nodes

Making Constraints “Soft”

- Want to minimize

$$\min_{\beta} F(\beta) = \sum_{d \in D, l \in L} h(p_l - p_d) + \lambda \|\beta\|^2$$

- Loss function: $h(x) = 0$ if $x < 0$, else $h(x) = x^2$



Regularization parameter

Penalty for violating the constraint that $p_d > p_l$

Solving the Problem: Intuition

- How to minimize F ?

$$\min_{\beta} F(\beta) = \sum_{d \in D, l \in L} h(p_l - p_d) + \lambda \|\beta\|^2$$

- Both p_l and p_d depend on β
 - Given β assign edge weights $a_{uv} = f_\beta(u, v)$
 - Using $Q=[a_{uv}]$ (edge strengths), compute PageRank scores p
 - Rank nodes by decreasing score
- **Goal:** want to find β such that $p_l < p_d$

Solving the Problem: Intuition

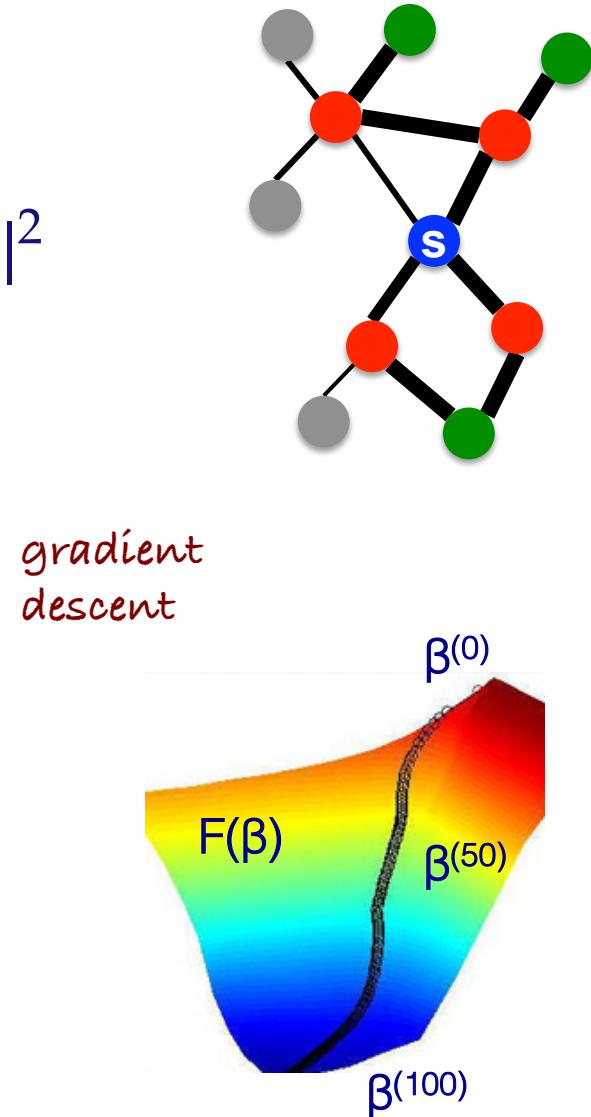
- How to minimize $F(\beta)$?

$$\min_{\beta} F(\beta) = \sum_{d \in D, l \in L} h(p_l - p_d) + \lambda \|\beta\|^2$$

- Idea
 - Start with some random $\beta^{(0)}$
 - Evaluate the derivative of $F(\beta)$ and do a small step in the opposite direction

$$\beta^{(t+1)} = \beta^{(t)} - \eta \frac{\partial F(\beta^{(t)})}{\partial \beta}$$

- Repeat until convergence



Data: Facebook

- Facebook Iceland network
 - 174,000 nodes (55% of population)
 - Avg. degree 168
 - Avg. person added 26 friends/month
- For every node **s**:
 - Positive examples
 - $D = \{\text{New friendships created in Nov '09}\}$
 - Negative examples
 - $L = \{\text{Other nodes } s \text{ that did not create new links to}\}$
 - Limit to friends of friends (FoFs) – 2-hop neighborhood
 - On avg. there are 20,000 FoFs (maximum is 2 million)

Experimental Setting

- Node and Edge features for learning
 - Node: Age, Gender, Degree
 - Edge: Age of an edge, Communication, Profile visits, Co-tagged photos
- Evaluation
 - Precision at top 20
 - We produce a list of 20 candidates
 - By taking top 20 nodes x with highest PageRank score p_x
 - Measure to what fraction of these nodes s actually links to

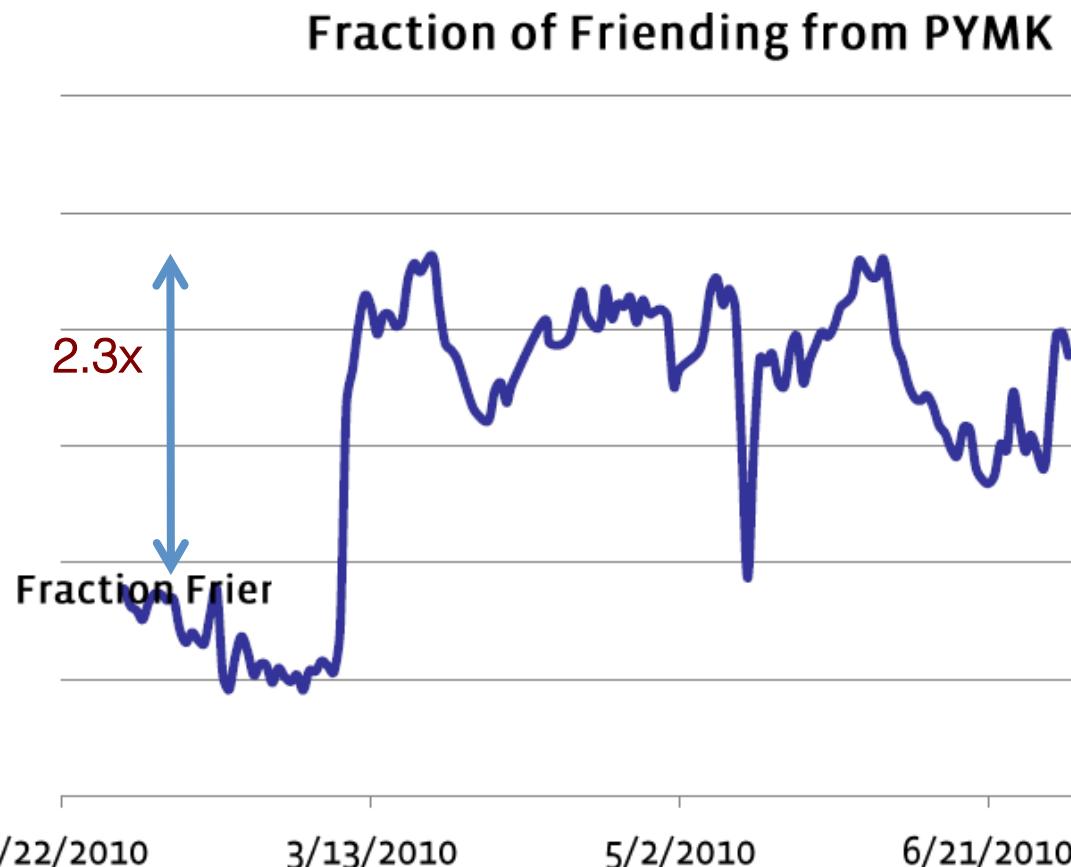
Experimental Results – Facebook (1/2)

- **Facebook:** Predict future friends
 - Adamic/Adar index already works great
 - Supervised Random Walks (SRW) gives slight improvement

Learning Method	Prec@Top20	Recommend 20 potential friendships in FB
Random Walk with Restart	6.80	
Adamic-Adar	7.35	
Common Friends	7.35	
Degree	3.25	
SRW: one edge type	6.87	
SRW: multiple edge types	7.57	

Experimental Results – Facebook (2/2)

- 2.3x improvement over previous FB-PYMK (People You May Know)



Experimental Results - Co-Authorship

- Arxiv Hep-Ph collaboration network
 - Poor performance of unsupervised methods
 - SRW gives a boost of 25%

Learning Method	Prec@Top20
Random Walk with Restart	3.41
Adamic-Adar	3.13
Common Friends	3.11
Degree	3.05
SRW: one edge type	4.24
SRW: multiple edge types	4.25

Link Prediction - Summary

- Node proximity measures
 - Unsupervised and supervised
- Supervised random walks
 - Experiments on Facebook data
- WTF: The “who to follow” service at Twitter
 - Combination of random walks with the SALSA algorithm (a variant of HITS)

Thank You!

