

Network Science Analytics

Option Applied Math and M.Sc. in DSBA

Lecture 5B
Graph similarity

Fragkiskos Malliaros

Friday, March 2, 2018

Acknowledgements

- The lecture is partially based on material by
 - Jure Leskovec, Stanford University
 - Aaron Clauset, CU Boulder
 - Manos Papaggelis, York University
 - Gonzalo Mateos, University of Rochester
 - Albert-László Barabási, Northeastern University
 - Christos Faloutsos, CMU
 - Panagiotis Tsaparas, University of Ioannina
 - Danai Koutra, University of Michigan
 - Michalis Vazirgiannis, Ecole Polytechnique
 - Karsten M. Borgwardt, University of Munich
 - R. Zafarani, M. A. Abbasi, and H. Liu, Social Media Mining: An Introduction, Cambridge University Press, 2014. Free book and slides at <http://socialmediamining.info/>

Thank you!

In This Lecture

- Graph similarity
 - Problem definition and applications
 - Known node correspondence
 - Unknown node correspondence

Graph similarity

Graph Similarity

Given two graphs \mathbf{G}_1 and \mathbf{G}_2 from the space of graphs \mathcal{G} the problem of graph similarity is to find a mapping

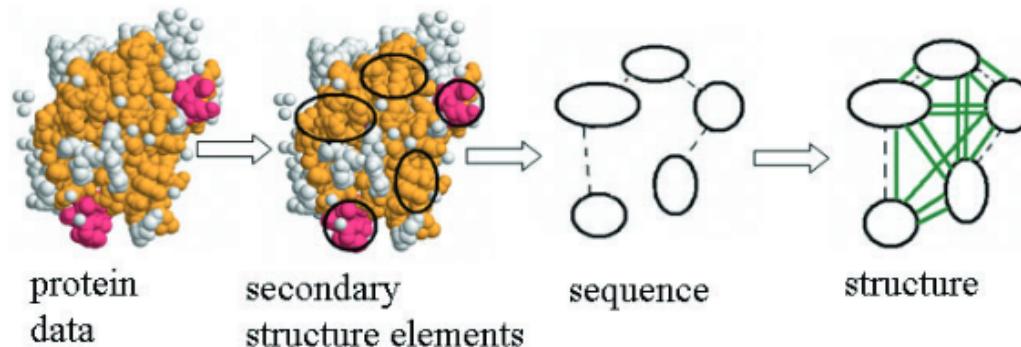
$$s : \mathcal{G} \times \mathcal{G} \rightarrow \mathbb{R}$$

such that $s(\mathbf{G}_1, \mathbf{G}_2)$ quantifies the similarity of \mathbf{G}_1 and \mathbf{G}_2

- Very important problem with numerous applications
 - Recall that data is often represented as graphs
 - Most machine learning algorithms make decisions based on the similarity or distance between pairs of data points (i.e., instances)

Applications – Protein Function Prediction

- For each protein, create a graph that contains information about its
 - Structure
 - Sequence
 - Chemical properties

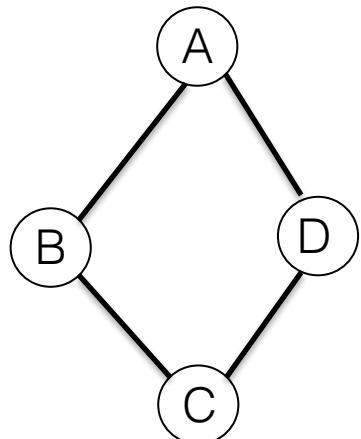


- Use graph similarity methods
 - To measure the structural similarity between proteins
 - To predict the function of proteins

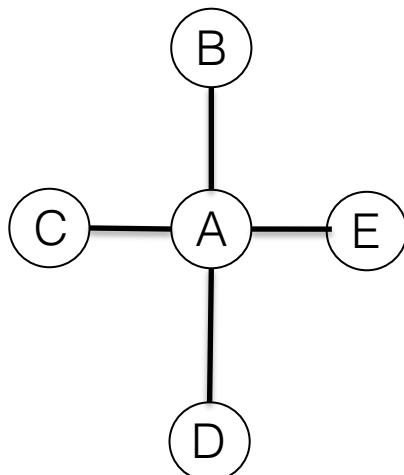
Applications – Classification of Molecules

Dataset of **known** molecules

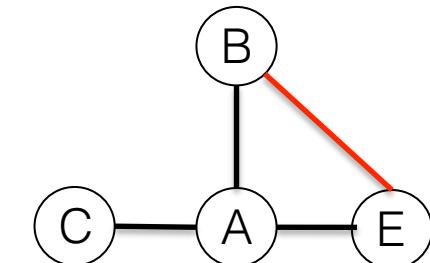
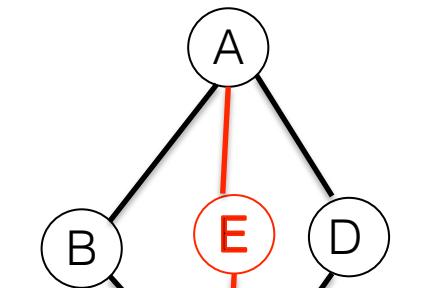
Toxic



Non-toxic



Unknown molecules



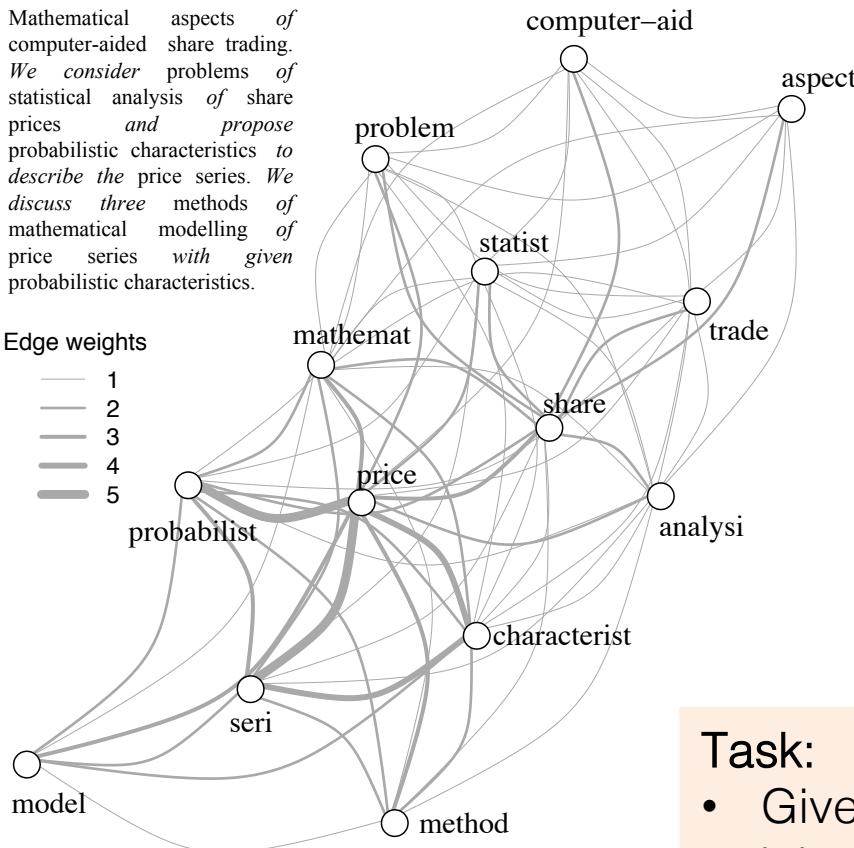
Task:

- Given a set of molecules that are either toxic or non-toxic
- Predict** the class of unknown molecules

Applications – Text Categorization

Mathematical aspects of computer-aided share trading. We consider problems of statistical analysis of share prices and propose probabilistic characteristics to describe the price series. We discuss three methods of mathematical modelling of price series with given probabilistic characteristics.

Edge weights



Represent a **document** as a graph

- **Nodes:** terms of the document
- **Edges:** co-occurrence within a fixed-size sliding window

Task:

- Given a set of documents and their class labels (e.g., positive/negative reviews)
- Correctly classify documents

Applications – Malware Detection

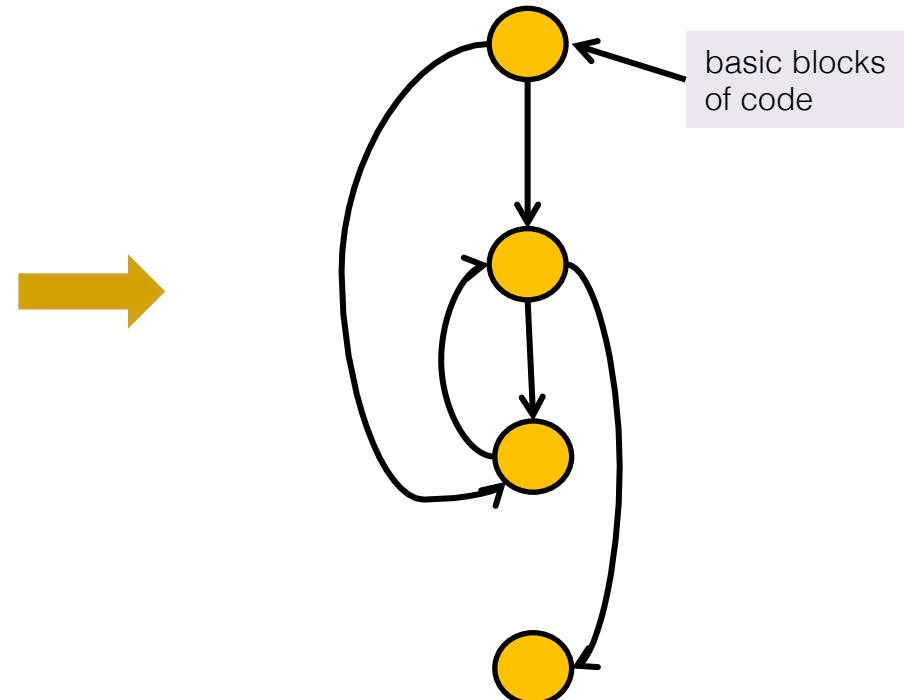
Computer program

```
# a scheduler class, to schedule and run events after a delay
class Scheduler:
    def __init__(self):
        # begin with no events
        self.events = []

    # after the delay, run the function
    def schedule(self, delay, function):
        if delay <= 0:
            # if no delay, run function straight away
            function()
        else:
            # queue the function
            self.events.append([delay, function])
            # sort the queue, smallest delay first
            self.events.sort(key = lambda event: event[0])

    def run(self):
        # loop while there are events to process
        while len(self.events) > 0:
            # get the first event
            delay, function = self.events.pop(0)
            # subtract time from remaining events
            for event in self.events:
                event[0] -= delay
            # wait for required delay
            pyb.delay(delay)
            # run function
            function()
```

Control flow graph



- Compare the control flow graph of the program against a set of control flow graphs of known malware
- If it contains a similar subgraph to these graphs, then there is probably malicious code running in the program

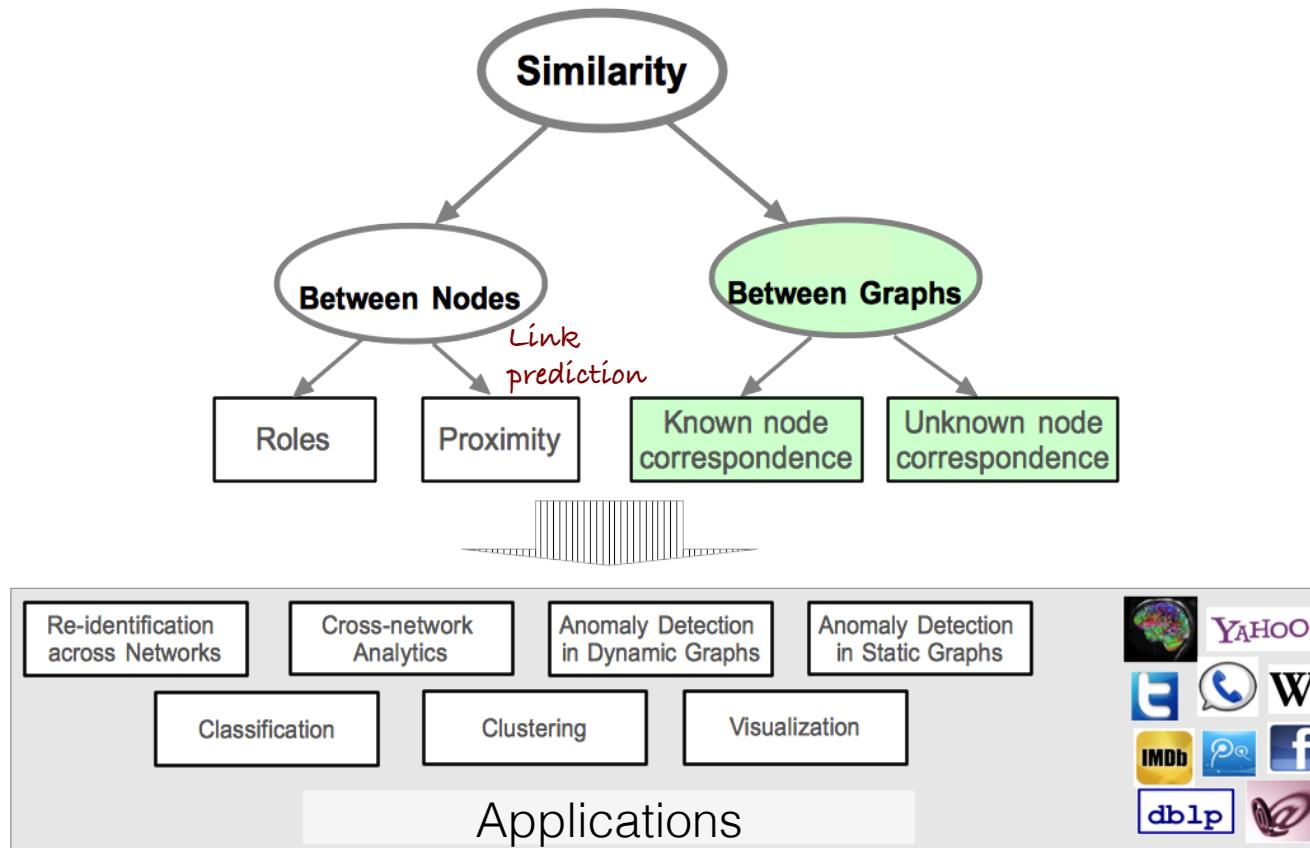
More Applications

- Bioinformatics
- Chemo-informatics
- Comparison of electric circuits
- Anomaly detection
- Relational entity disambiguation
- XML document similarity
- Entity classification
- Object tracking
- Process mining in the business domain
- ...

It's Not a Trivial Task

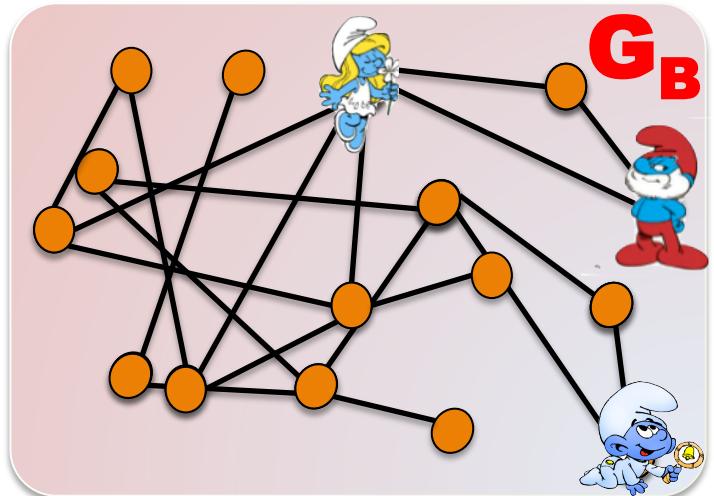
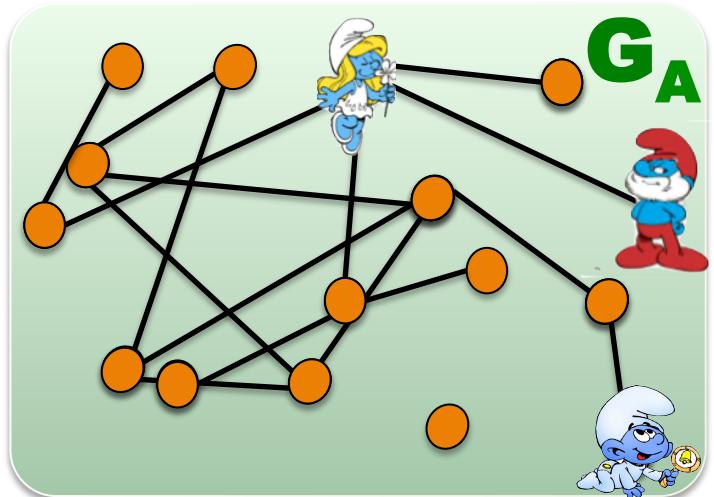
- Although graph similarity (GS) seems to be a tractable problem, in fact it is a complex task
- Many problems related to GS are **computationally difficult**
 - Subgraph isomorphism [will be defined later]
 - Finding the largest common subgraph
- We are mainly interested in GS algorithms that run in polynomial time
 - However, in large scale networks even poly-time algorithms are not applicable (e.g., algorithms with running time $O(n^4)$)

Similarity in Graphs



Graph similarity: known node correspondence

Problem Definition: Graph Similarity



Given:

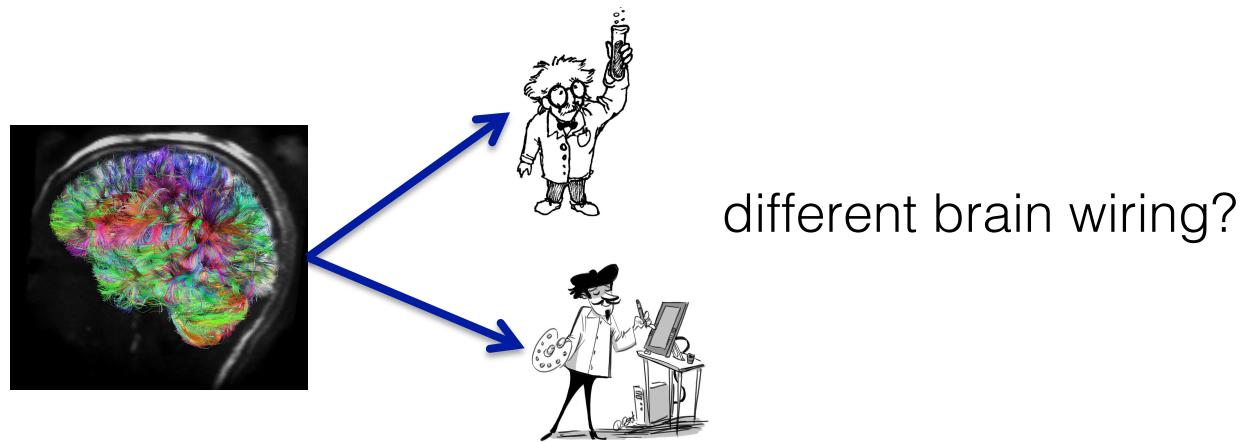
- (i) two graphs with the **same** nodes and **different** edge sets
- (ii) node correspondence

Find: similarity score

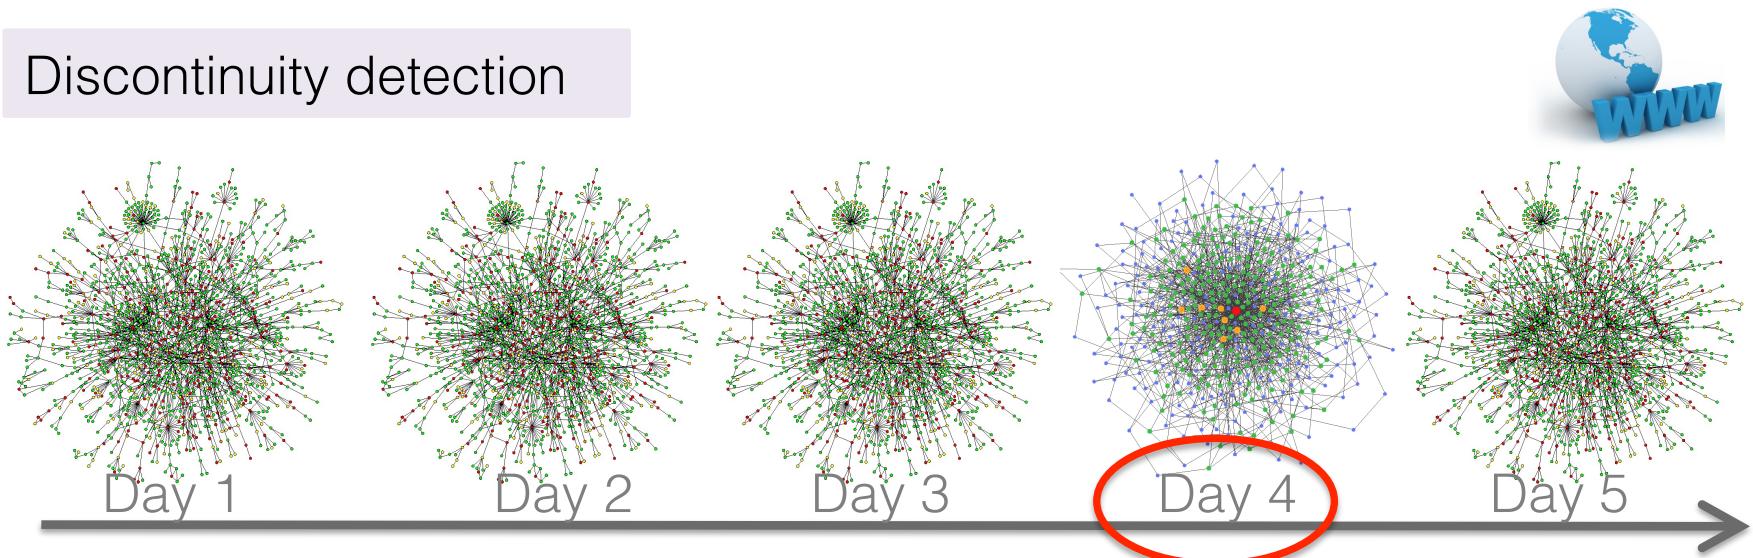
$$s \in [0, 1]$$

Applications (1/2)

Classification



Discontinuity detection



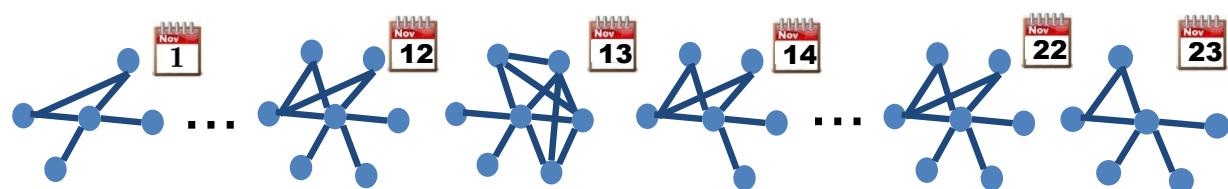
Applications (2/2)



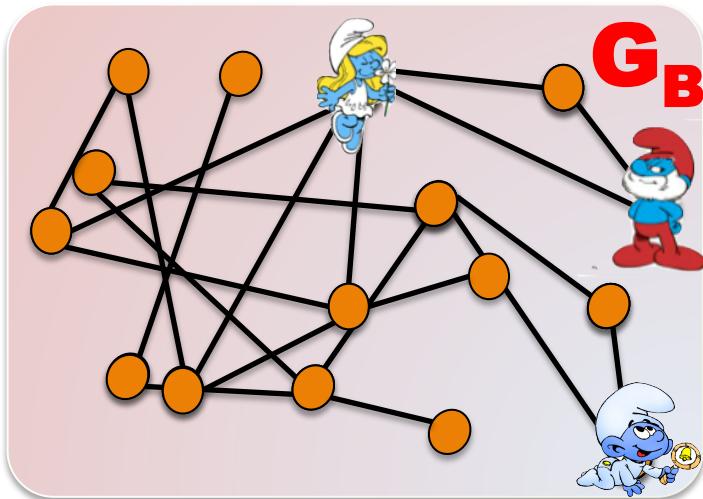
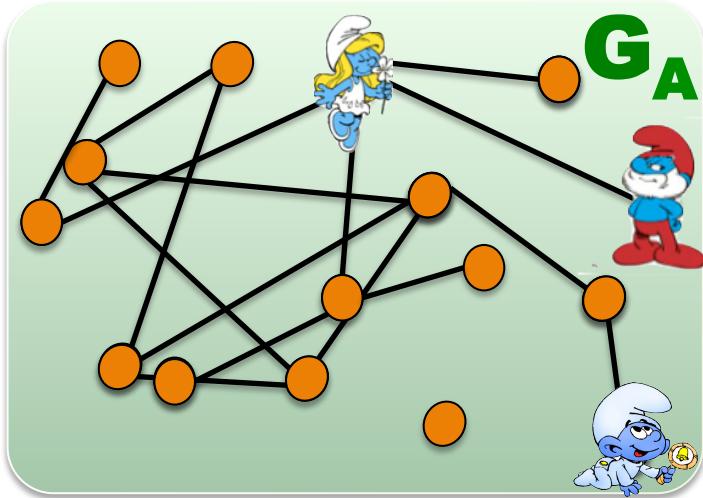
Behavioral Patterns

FB message graph vs. wall-to-wall network

Intrusion detection



A First Solution



Edge Overlap (EO)

of common edges
(normalized or not)

Is it a good measure?

Edge Overlap – Example

$$\text{EO}(B10, mB10) == \text{EO}(B10, mmB10)$$



Barbell graph

G_A

B10

G_B

mB10

G_A

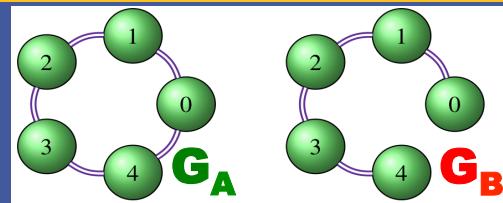
B10

G_{B'}

mmB10

Different impact in information flow

Similar If ...



1. "... they share many vertices and/or edges"

$$VEO = \frac{|E_A \cap E_B| + |V_A \cap V_B|}{|E_A| + |E_B| + |V_A| + |V_B|}$$

Vertex/Edge Overlap:
 $O(|V_A| + |V_B| + |E_A| + |E_B|)$

2. "... the rankings of their vertices are similar"

VR = rank correlation of node PageRank

Vertex Ranking:
 $O(|V_A| + |V_B|)$

3. "... their edge weights are similar"

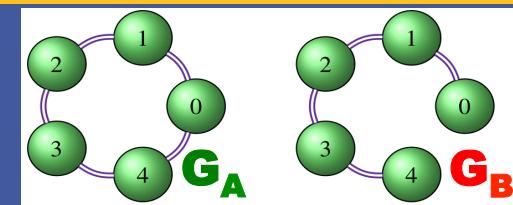
$$d(G_A, G_B) = \frac{1}{|E_A \cup E_B|} \sum_e \frac{|w_{GA}(e) - w_{GB}(e)|}{\max\{w_{GA}(e), w_{GB}(e)\}}$$

Weighted distance:
 $O(|E_A| + |E_B|)$

[Papadimitriou, Dasdan, Garcia-Molina '10; Bunke '06,
Shoubridge et al. '02, Dickinson et al. '04]

Similar If ...

4. "... they have similar subgraphs"



Maximum Common Subgraph (MCS)
(NP-complete problem)

vertex MCS Distance

$$d(\mathbf{G}_A, \mathbf{G}_B) = 1 - \frac{|\text{mcs}(\mathbf{V}_A, \mathbf{V}_B)|}{\max\{|\mathbf{V}_A|, |\mathbf{V}_B|\}}$$

Edge MCS Distance

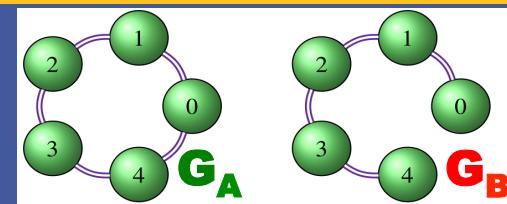
$$d(\mathbf{G}_A, \mathbf{G}_B) = 1 - \frac{|\text{mcs}(\mathbf{E}_A, \mathbf{E}_B)|}{\max\{|\mathbf{E}_A|, |\mathbf{E}_B|\}}$$

5. "... if we need few node/edge additions/deletions
to transform \mathbf{G}_A to \mathbf{G}_B "

Graph Edit Distance

$$\text{GED}(\mathbf{G}_A, \mathbf{G}_B) = |\mathbf{V}_A| + |\mathbf{V}_B| - 2|\mathbf{V}_A \cap \mathbf{V}_B| + |\mathbf{E}_A| + |\mathbf{E}_B| - 2|\mathbf{E}_A \cap \mathbf{E}_B|$$

Similar If ...



6. "... they have similar fingerprints"

Signature similarity

b-bit fingerprint of \mathbf{G}_A :

1	0	1	0	1
---	---	---	---	---

b-bit fingerprint of \mathbf{G}_B :

0	0	1	0	1
---	---	---	---	---

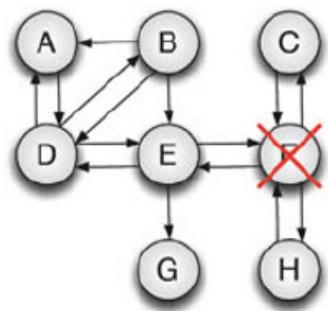


Hamming distance: 1

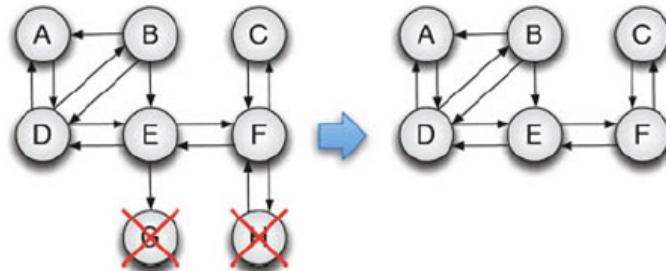
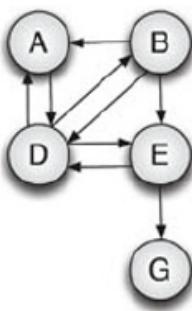
The fingerprint is created by applying a hash function on the nodes and edges (features):

- Nodes: use the PageRank score
- Edges (u, v) : The PageRank score of u multiplied by the out-degree of u

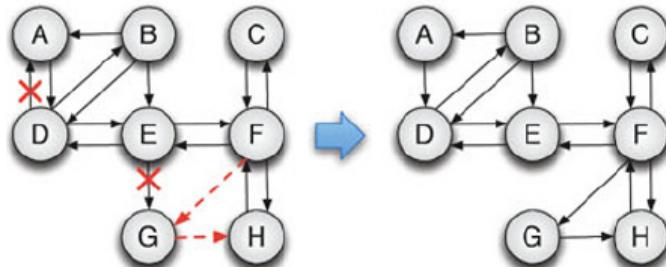
Application: Web Graph Anomaly Detection



Anomaly I : Missing connected subgraph



Anomaly II: Missing random nodes



Anomaly III: Connectivity change
(same # of nodes/edges but different
connectivity for some nodes)

Anomaly detection in the Web graph:

- Given a sequence of web graphs G_1, \dots, G_n
- Compute similarity between G_i and G_{i+1}
- Construct time series based on similarity scores
- Detect anomalies

	Missing Subgraph	Missing Vertices	Topological Changes
Vertex/Edge Overlap	✓✓	✓	✗
Vector Similarity	✓✓	✓	✓
Vertex Ranking	✓✓	✗	✗
Sequence Similarity	✓	✗	✓✓
Signature Similarity	✓✓	✓✓	✓✓

[Papadimitriou, Dasdan, Garcia-Molina '10]

Overview of Graph Similarity Functions

- Many similarity functions can be defined
- What **properties** should a **good** similarity function have?

Axioms of Graph Similarity Functions

A1. Identity property

$$\text{sim}(\text{graph A}, \text{graph A}) = 1$$

A2. Symmetric property

$$\text{sim}(\text{graph A}, \text{graph B}) = \text{sim}(\text{graph B}, \text{graph A})$$

A3. Zero property

$$\text{sim}(\text{graph A}, \text{graph C}) = 0$$

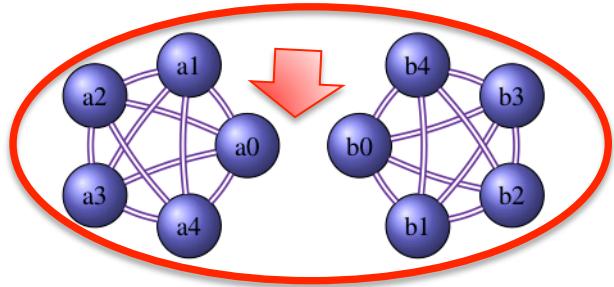
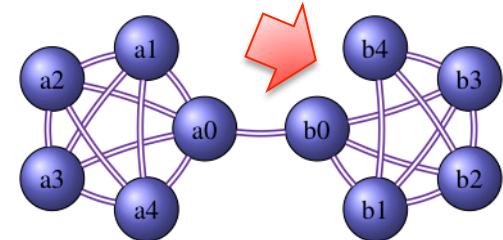
Desired Properties

- Intuitiveness
 - P1. Edge Importance
 - P2. Weight Awareness
 - P3. Edge-“Submodularity”
 - P4. Focus Awareness
- Scalability

Desired Properties

- Intuitiveness

- P1. Edge Importance
 - P2. Weight Awareness
 - P3. Edge-“Submodularity”
 - P4. Focus Awareness



- Scalability

Creation of disconnected components matters more than small connectivity changes

Desired Properties

- Intuitiveness

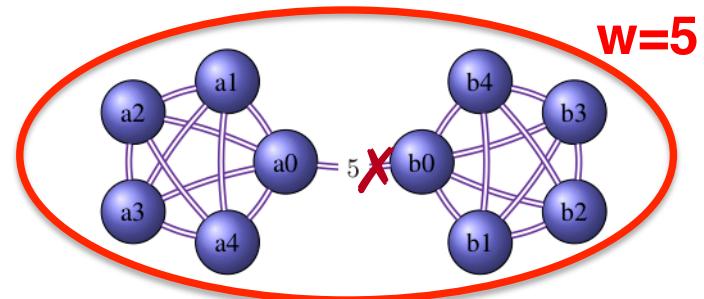
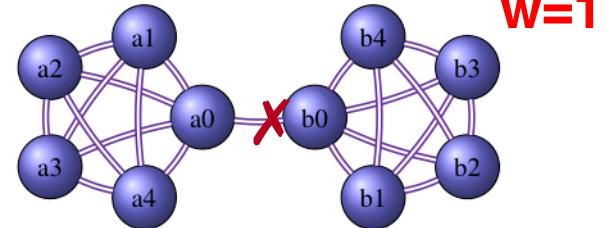
P1. Edge Importance

P2. Weight Awareness

P3. Edge-“Submodularity”

P4. Focus Awareness

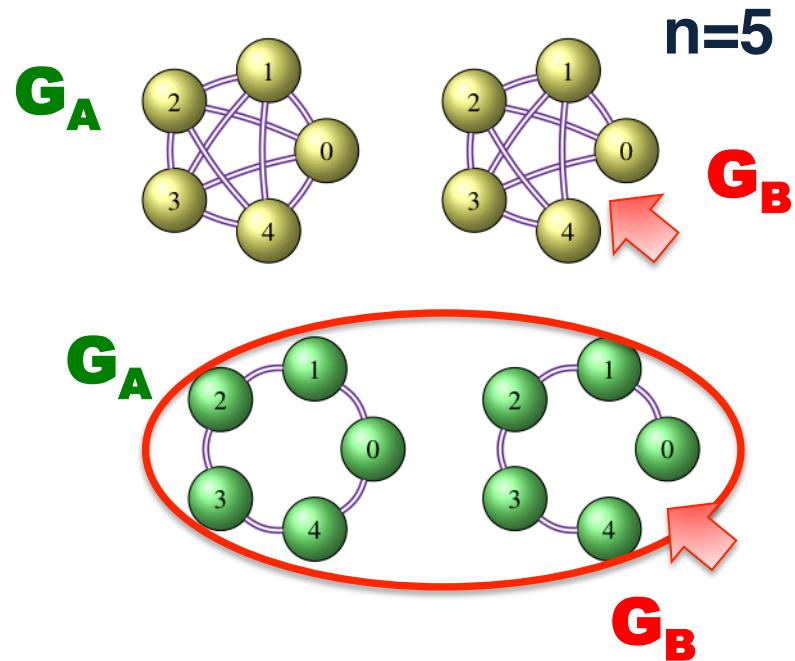
- Scalability



The bigger the edge weight, the more the edge change matters

Desired Properties

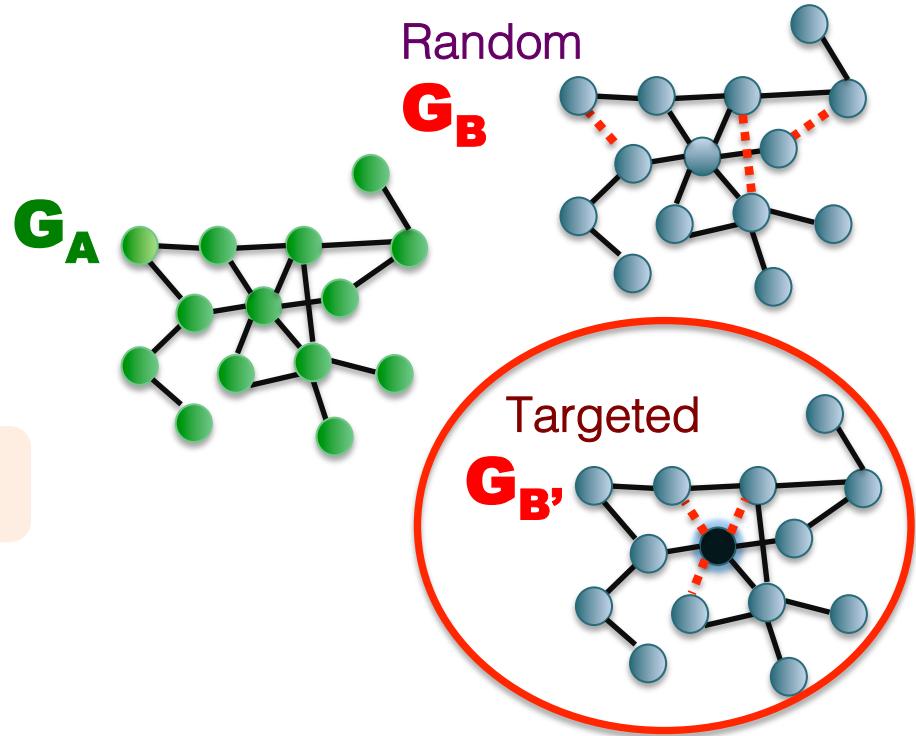
- Intuitiveness
 - P1. Edge Importance
 - P2. Weight Awareness
 - P3. Edge-“Submodularity”
 - P4. Focus Awareness
- Scalability



The sparser the graphs, the more important is a “fixed” change (e.g., removal of an edge)

Desired Properties

- Intuitiveness
 - P1. Edge Importance
 - P2. Weight Awareness
 - P3. Edge-“Submodularity”
 - P4. Focus Awareness
- Scalability



Targeted changes are more important than **random** changes of the same extend

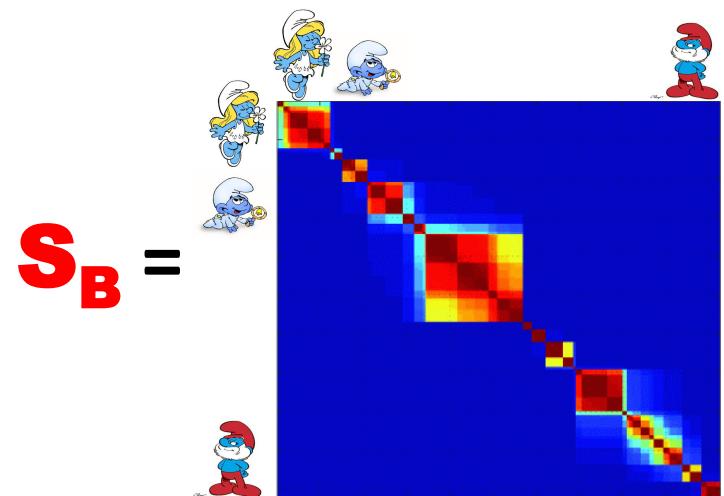
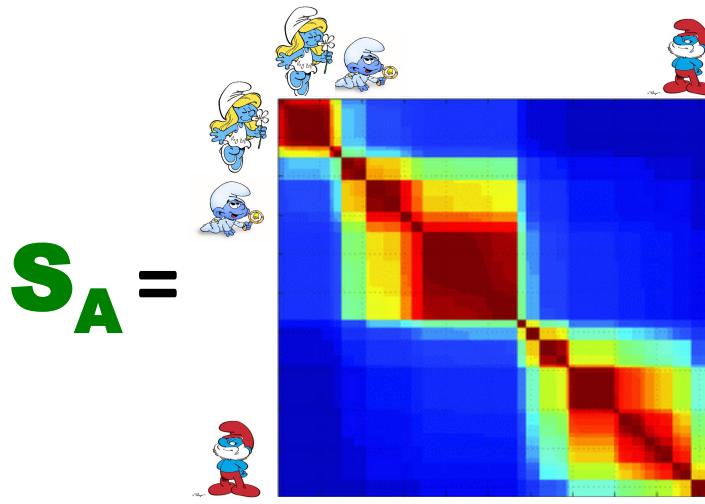
Properties of State-of-the-Art Methods

Metric	importance	weight	sparsity	focus
	P1	P2	P3	P4
Vertex/Edge Overlap	✗	✗	✗	?
Graph Edit Distance (XOR)	✗	✗	✗	?
Signature Similarity	✗	✓	✗	?

Looking for similarity functions that satisfy these properties

DELTACon

1. Find the pairwise node influence, S_A and S_B
 - Node influence = node affinity
2. Find the similarity between S_A and S_B



[Koutra, Faloutsos, Vogelstein '13]

How to Measure Node Affinity

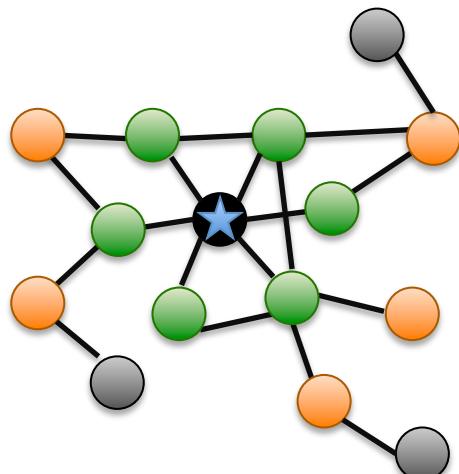
- Any random walk-based method can be used
 - E.g., PageRank, Random walk with restarts (RWR), personalized PR
 - Use Fast Belief Propagation (FaBP) *similar to Personalized RWR*
- Attenuating Neighboring Influence for small ε :

$$S = [I + \varepsilon^2 D - \varepsilon A]^{-1} \approx$$

S : similarity matrix

$$\approx [I - \varepsilon A]^{-1} = I + \varepsilon A + \varepsilon^2 A^2 + \dots$$

1-hop *2-hops* ...
↓ ↓



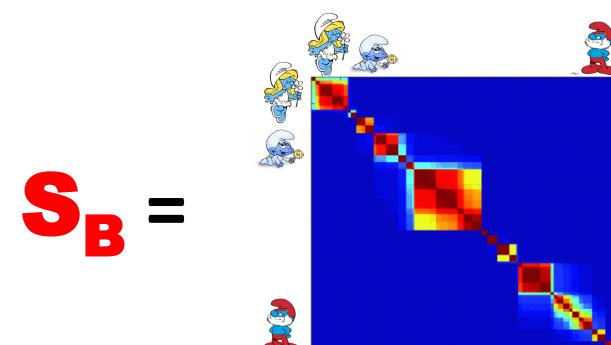
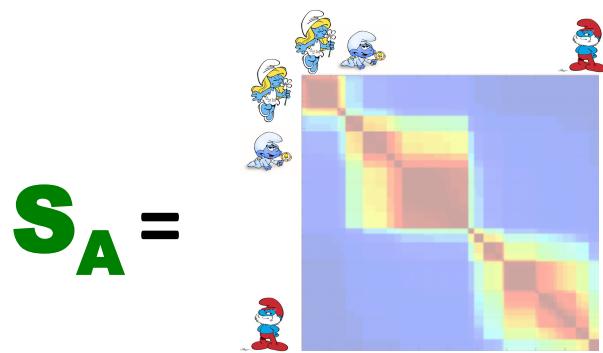
Why FaBP?

- Sound theoretical background
- Fast
- Intuitive: consider k -hop away information
 - Differences in long paths have smaller effect in the similarity

Graph Similarity with DELTA CON

1. Find the pairwise node influence, S_A and S_B
 - Node influence = node affinity
2. Find the similarity between S_A and S_B

$$sim(\mathbf{S}_A, \mathbf{S}_B) = \frac{1}{1 + \sqrt{Euclidean\ Dist.}} = \frac{1}{1 + \sqrt{\sum_{i,j} (\sqrt{s_{Ai,j}} - \sqrt{s_{Bi,j}})^2}}$$



$$s(\mathbf{S}_A, \mathbf{S}_B) = 0.3$$

[Koutra, Faloutsos, Vogelstein '13]

DELTACon - The Algorithm

Input: graphs $G_1 = (V, E_1)$ and $G_2 = (V, E_2)$

Output: similarity $s(G_1, G_2)$

$$1. \quad S_1 = [I + \epsilon^2 D_1 - \epsilon A_1]^{-1}$$

$S_{1,ij}$: affinity/influence of node i to node j in G_1

$$2. \quad S_2 = [I + \epsilon^2 D_2 - \epsilon A_2]^{-1}$$

$$3. \quad d(G_1, G_2) = \text{Root Euclidean Distance}(S_1, S_2)$$

$$4. \quad s(G_1, G_2) = \frac{1}{1 + d}$$

- Quadratic complexity (n^2 affinity scores need to be computed)
- DELTACon₀: faster, linear approximation algorithm
 - Split the node set into g groups and compute $n \times g$ affinity scores

Code: <http://www.cs.cmu.edu/~dkoutra/CODE/deltacon.zip>

Properties of Similarity Methods

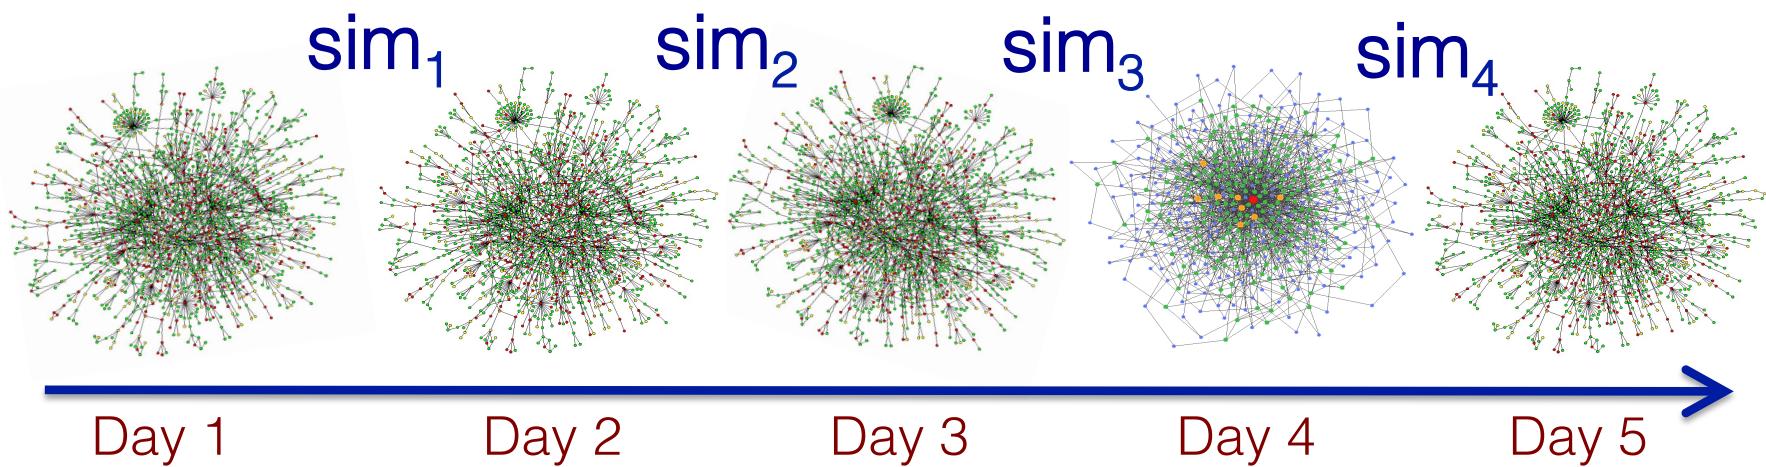
Metric	importance	weight	sparsity	focus
	P1	P2	P3	P4
Vertex/Edge Overlap	✗	✗	✗	?
Graph Edit Distance (XOR)	✗	✗	✗	?
Signature Similarity	✗	✓	✗	?
DELTACon	✓	✓	✓	✓
DELTACon ₀	✓	✓	✓	✓

[Koutra, Faloutsos, Vogelstein '13]

Temporal Anomaly Detection (1/2)

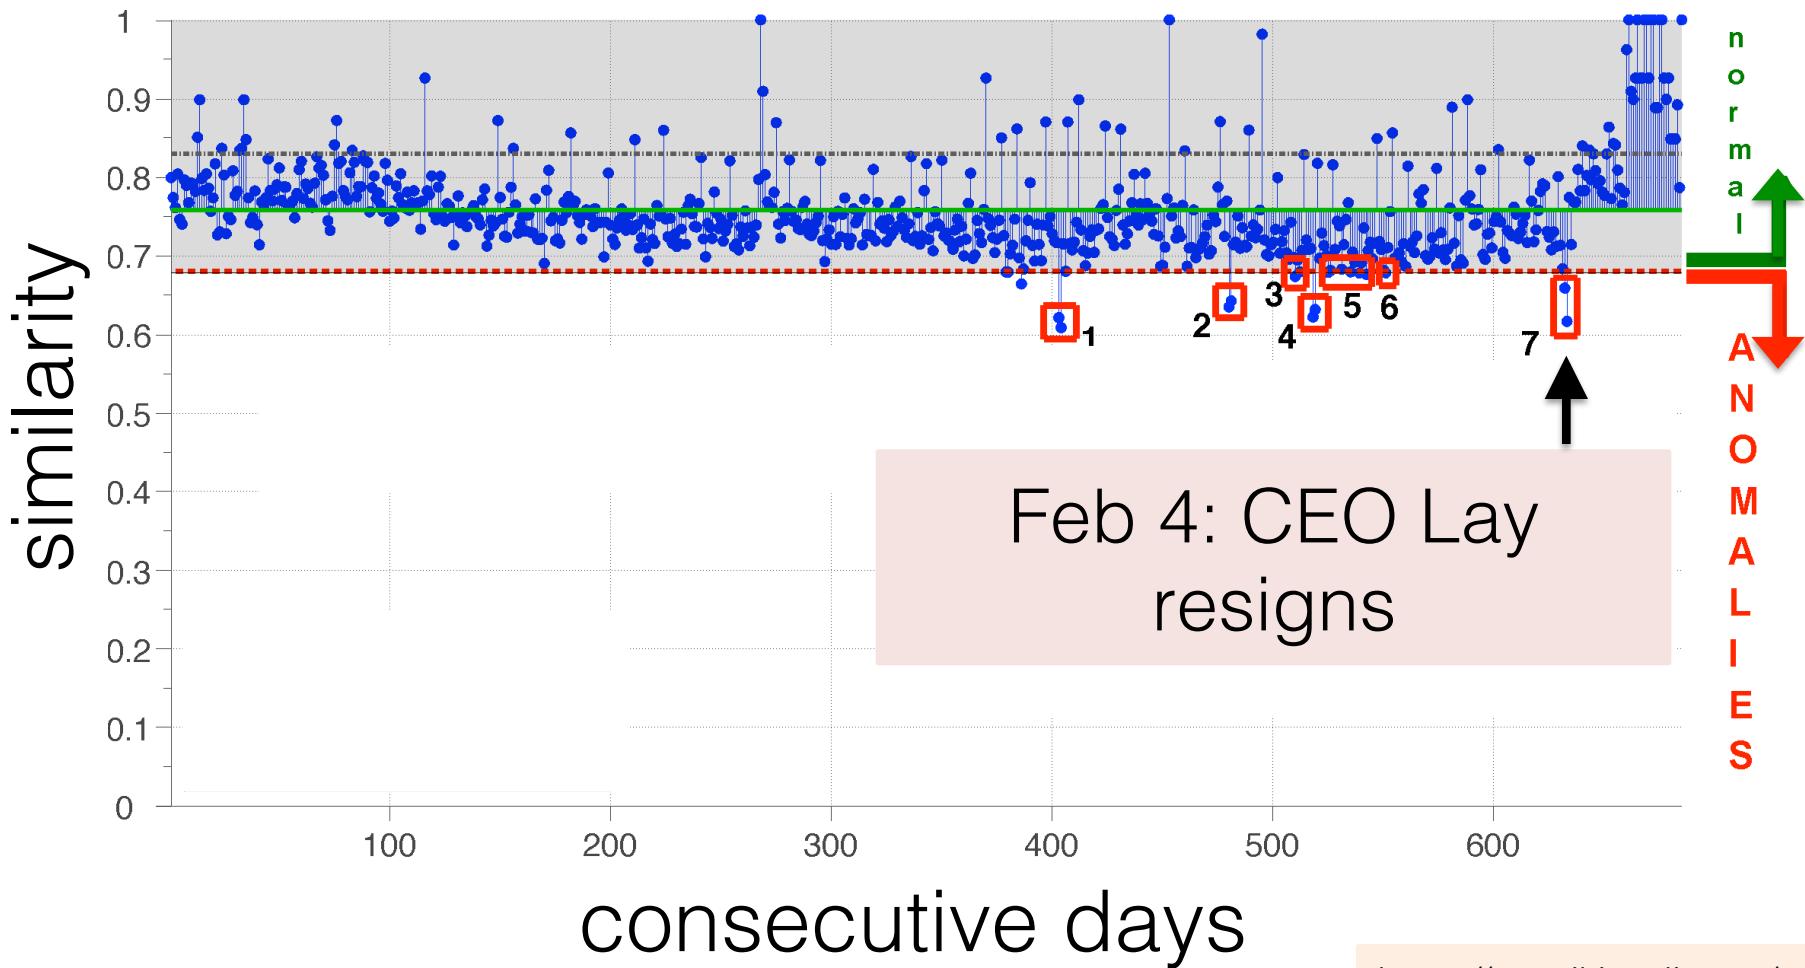
Nodes: employees

Edges: email exchange



[Koutra, Faloutsos, Vogelstein '13]

Temporal Anomaly Detection (2/2)



Use the method of Quality Control with Individual Moving Range to define the threshold below which days are anomalous (red line): $\text{median} \pm 3\sigma$

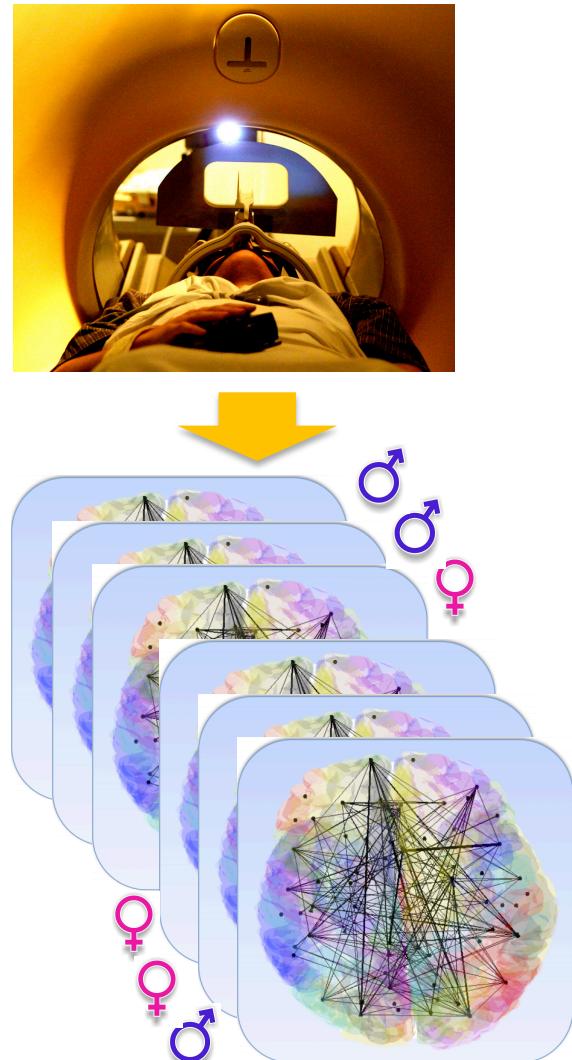
[https://en.wikipedia.org/
wiki/Enron_scandal](https://en.wikipedia.org/wiki/Enron_scandal)

Brain Connectivity Graph Clustering (1/2)

- 114 brain graphs (Multimodal MRI)
 - One graph per person
 - **Nodes:** 70 cortical regions
 - **Edges:** connections
- **Attributes:** gender, IQ, age...

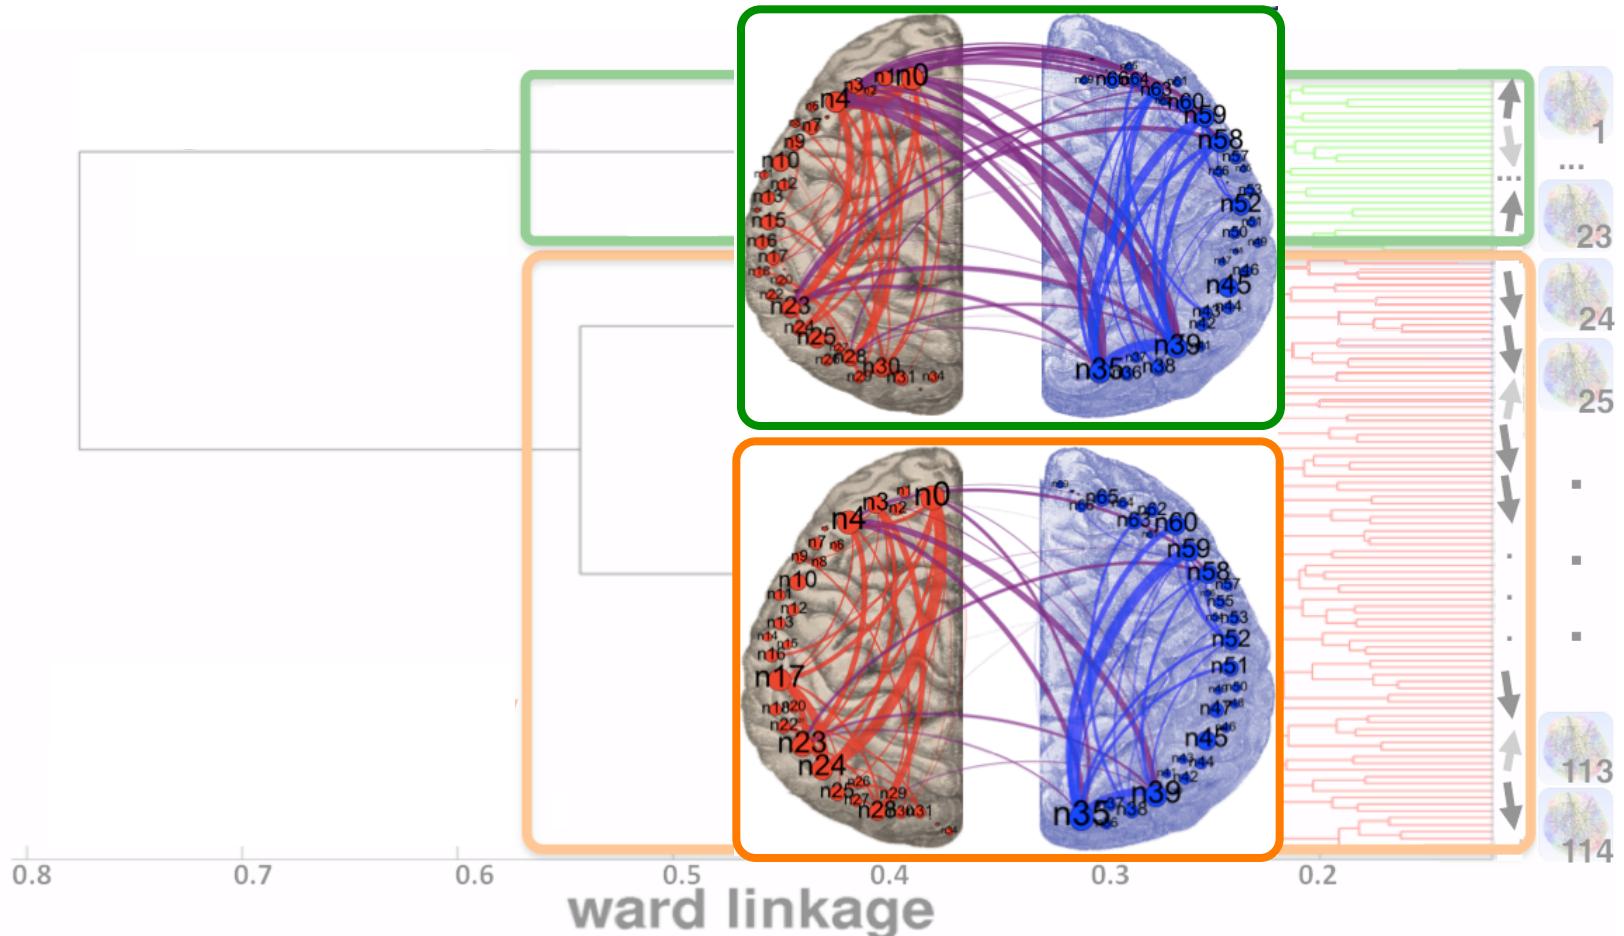
Experimental setup

- Compute similarities with **DELTACon**
- Perform hierarchical clustering using Ward's method (https://en.wikipedia.org/wiki/Ward%27s_method)
- Apply t-test on the attributes of the extracted groups to detect differences



[Koutra, Faloutsos, Vogelstein '13]

Brain Connectivity Graph Clustering (2/2)



Significant difference in t-test:
p-value = 0.0057

[Koutra, Faloutsos, Vogelstein '13]

CCI: Composite Creative Index
- Related to a person's performance on a series of creative tasks

Graph similarity: unknown node correspondence

Main Idea

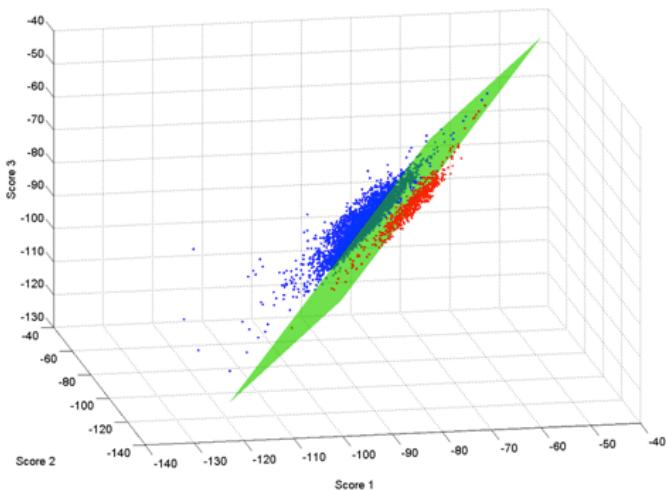
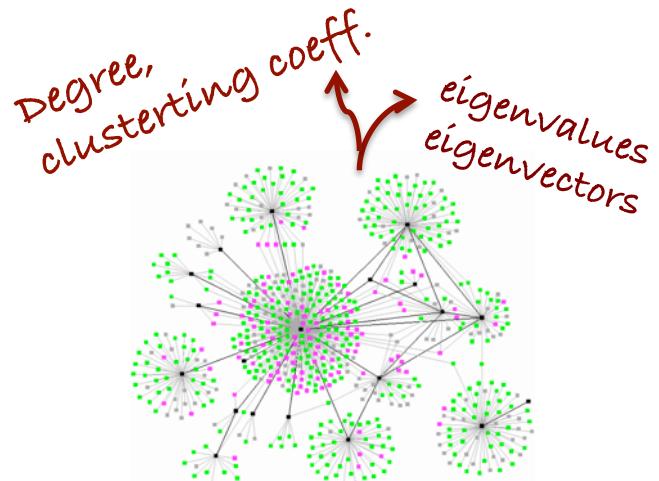
1. Extract graph features
2. $s(G_A, G_B) =$ “similarity” between the features

Or

apply some classifier

Which features?

- Consider spectral information



λ -distance

- Eigenvalues of G_A : $\lambda_{A1} \leq \dots \leq \lambda_{At}$
- Eigenvalues of G_B : $\lambda_{B1} \leq \dots \leq \lambda_{Bs}$

$$d(G_A, G_B) = \sqrt{\sum (\lambda_{Ai} - \lambda_{Bi})^2}$$

[Bunke '06, Wilson '08, ElGhawalby '08]

λ -distance: Variations

- Eigenvalues of different matrices

- Adjacency

$$A = \begin{matrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{matrix}$$

- Laplacian

$$0 = \lambda_{A1} \leq \dots \leq \lambda_{At}$$

$$L = D - A$$

$$d_1 d_2 d_3 = \begin{matrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{matrix}$$

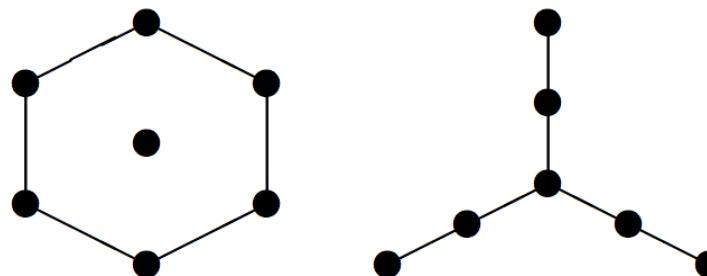
- Normalized Laplacian

$$0 = \lambda_{A1} \leq \dots \leq \lambda_{At} \leq 2$$

$$L_{\text{norm}} = \frac{D^{-1/2}}{d_1^{-0.5} d_2^{-0.5} d_3^{-0.5}} \begin{matrix} A \\ 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{matrix} \frac{D^{-1/2}}{d_1^{-0.5} d_2^{-0.5} d_3^{-0.5}}$$

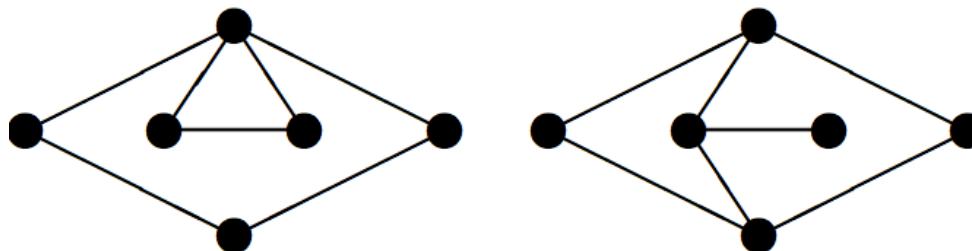
λ -distance: Disadvantages

- 1) co-spectral graphs with **different structure**



Non-isomorphic graphs, co-spectral w.r.t. \mathbf{A}

Non-isomorphic graphs, co-spectral w.r.t. \mathbf{L}



[Haemers+ '04, Brouwer '09]

λ -distance: Disadvantages

- 1) co-spectral graphs with **different** structure
- 2) subtle changes in the graphs =>
big differences in the spectrum
- 3) $O(n^3)$ runtime
 - E.g., using SVD
- 4) Labeled graphs: nodes can be associated with labels
 - How can you take into account this information?
 - Use graph kernels



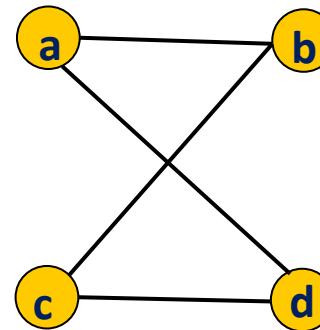
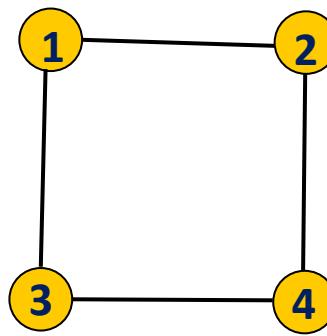
Graph Kernels

Graph Isomorphism

- Two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$ are isomorphic if there is a bijective mapping (one to one correspondence) $f: V_1 \rightarrow V_2$ such that $(v_i, v_j) \in E_1$ iff $(f(v_1), f(v_2)) \in E_2$
- The **Graph Isomorphism (GI)** is an “extreme” case of the graph similarity problem
 - It asks if two graphs are structurally equivalent to each other
 - No polynomial algorithm for GI is known
 - Quasipolynomial time $2^{(\log n)^{O(1)}}$ by L. Babai [Babai ‘15]

Example

The following graphs are isomorphic



Function f :

- $f(1) = a$
- $f(2) = b$
- $f(3) = d$
- $f(4) = c$

One-to-one correspondence between
the nodes of the two graphs

Subgraph Isomorphism

- The **Subgraph Isomorphism** problem asks if there is a subset of edges and nodes of G_1 that is isomorphic to a smaller graph G_2
 - NP-complete problem
- Practical problems
 - The runtime may grow exponentially
 - Can we have a polynomial time similarity measure?

Graph Kernels – The Idea

- Compare substructures of graphs that are computable in polynomial time
 - Walks
 - Shortest paths
 - Cyclic patterns
 - Trees
 - Graphlets

Our goal: Learning on graph data

- Represent a graph with a fixed-length feature vector
- Typical representation required by most of ML algorithms
- There is no straightforward way to transform graphs to such a representation

Kernel Methods in ML

- The function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a kernel if it is:
 - Symmetric $k(x, y) = k(y, x)$
 - Positive semi-definite $\forall x_1, x_2, \dots, x_n \in \mathcal{X}$, the matrix $\mathbf{K}_{ij} = k(x_i, x_j)$ is positive semi-definite
- If a function satisfies the above two conditions on a set \mathcal{X} , it is known that there exists a mapping $\phi : \mathcal{X} \rightarrow \mathbb{H}$ into a Hilbert space \mathbb{H} such that:

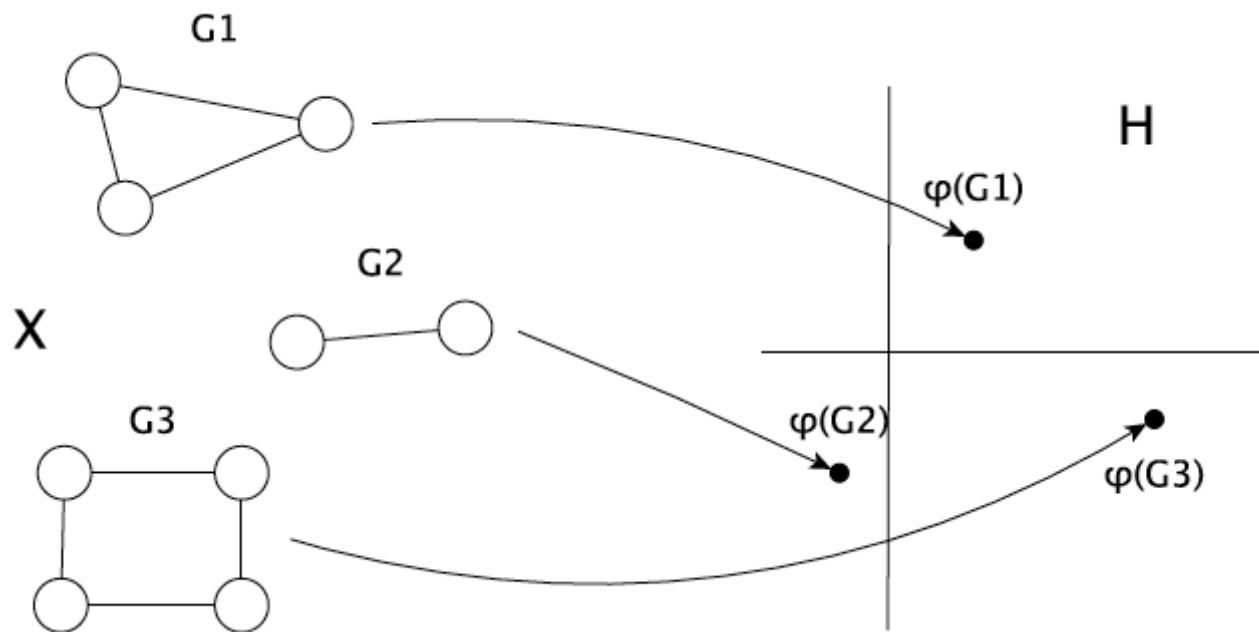
$$k(x, y) = \langle \phi(x), \phi(y) \rangle$$

for all $(x, y) \in \mathcal{X}^2$ where $\langle \cdot, \cdot \rangle$ is the inner product in \mathbb{H}

- Informally, $k(x, y)$ is a measure of similarity between x and y

Graph Kernels

- A graph kernel $k : \mathcal{G} \times \mathcal{G} \rightarrow \mathbb{R}$ is a kernel function over a set of graphs \mathcal{G}



Makes the family of kernel methods applicable to graphs

Kernel Trick in ML

- Many machine learning algorithms can be expressed only in terms of **inner products** between vectors
- Computing the explicit mappings $\Phi(x_1)$, $\Phi(x_2)$ and their inner product $\langle \Phi(x_1), \Phi(x_2) \rangle$ for the pair of graphs can be computationally expensive
- **Kernel trick:** avoid the explicit mapping by directly computing the inner product $\langle \Phi(x), \Phi(y) \rangle$ via the kernel function

Random Walk Kernel – The Idea

- **Idea:** count common walks in $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$
 - Walks are sequences of nodes that allow repetitions of nodes
 - Recall that, walks of length ℓ can be computed by taking the adjacency matrix A to the power of k
 - $A^\ell(i, j) = c$ means that c walks of length ℓ exist between nodes i and j
- How to find common walks in two graphs?
 - Use the **Product graph** G_x of G_1 and G_2

$$V_x = \{(v_1, v_2) : v_1 \in V_1, v_2 \in V_2\}$$

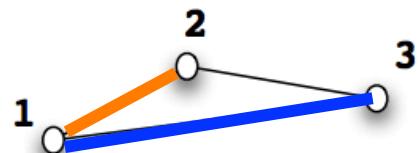
$$E_x = \{((v_1, v_2), (u_1, u_2)) : (v_1, u_1) \in E_1 \quad , (v_2, u_2) \in E_2\}$$

Nodes: pairs of vertices from G_1 and G_2

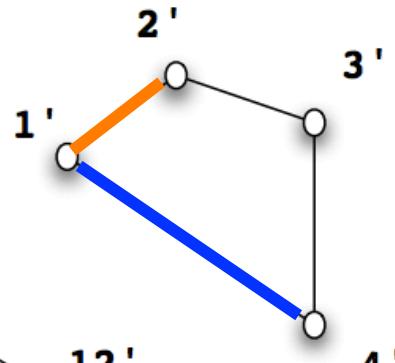
Draw edge if corresponding nodes of G_1 and G_2 are adjacent in G_1 and G_2

Product Graph

$$G_1 = (V_1, E_1)$$



$$G_2 = (V_2, E_2)$$



$$G_x = (V_x, E_x)$$

$$V_x = \{(v_1, v_2) : v_1 \in V_1, v_2 \in V_2\}$$

$$E_x = \{((v_1, v_2), (u_1, u_2)) : (v_1, u_1) \in E_1 \in V_1, (v_2, u_2) \in E_2\}$$

- Performing a random walk on G_x is equivalent to performing a simultaneous random walk on G_1 and G_2
- Common walks of length ℓ can be computed using A_x^ℓ

Random Walk Kernel – Computation

- Random walk kernel
 - Construct the product graph G_x
 - Count walks in this product graph
 - The random walk kernel can be computed as follows:

$$\begin{aligned} k_x(G_1, G_2) &= \sum_{i,j=1}^{|V_x|} \left[\sum_{\ell=0}^{\infty} \lambda^\ell \mathbf{A}_x^\ell \right]_{ij} \\ &= \mathbf{1}^T (\mathbf{I} - \lambda \mathbf{A}_x)^{-1} \mathbf{1} \end{aligned}$$

λ is a convergence factor

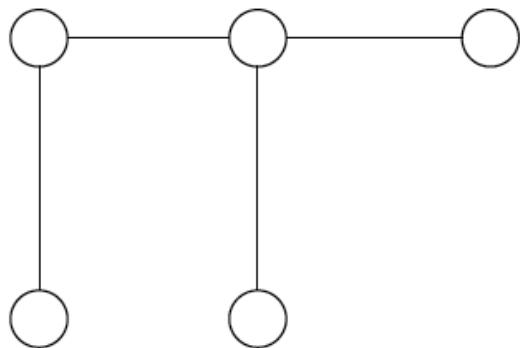
- If both graphs have size n (number of nodes)
 - $O(n^4)$ for computing the product graph (comparison for all pairs of edges)
 - Matrix multiplication or inversion of an $n^2 \times n^2$ matrix
 - Total complexity: $O(n^6)$
 - Can be reduced to $O(n^3)$ (using the Sylvester equation)

Shortest Path Kernels (1/3)

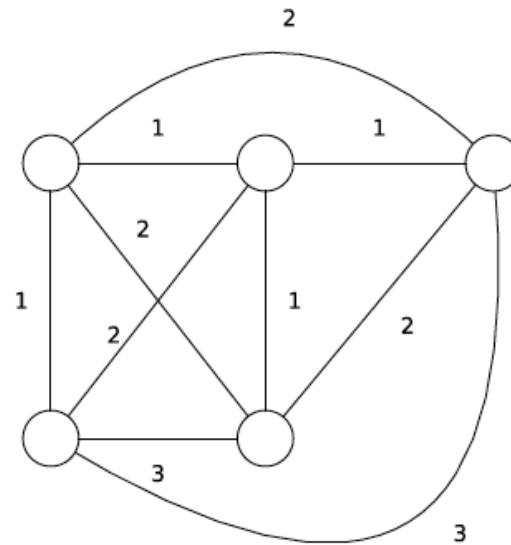
- The similarity of \mathbf{G}_1 and \mathbf{G}_2 is determined by the similarity of the length of shortest paths in the two graphs
- **Floyd transformation**
 - Compute the shortest-paths between all pairs of nodes of the input graph \mathbf{G} using the Floyd-Warshall algorithm
 - Create a shortest-path graph \mathbf{S} which contains the same set of nodes as the input graph \mathbf{G}
 - All nodes which are connected by a walk in \mathbf{G} are linked with an edge in \mathbf{S}
 - Each edge in \mathbf{S} is labeled by the shortest distance between its endpoints in \mathbf{G}

Shortest Path Kernels (2/3)

Floyd transformation



G



S

Shortest Path Kernels (3/3)

- Given the Floyd-transformed graphs $S_1 = (V_1, E_1)$ and $S_2 = (V_2, E_2)$ of G_1 and G_2 , the shortest path kernel is defined as

$$k(G_1, G_2) = \sum_{e_1 \in E_1} \sum_{e_2 \in E_2} k_{\text{walk}}^{(1)}(e_1, e_2)$$

where $k_{\text{walk}}^{(1)}$ is a kernel on edge walks of length 1

Complexity:
 $O(n^4)$

$$k_{\text{walk}}^{(1)} = \begin{cases} 1, & \text{if } \ell(e_1) = \ell(e_2) \\ 0, & \text{otherwise} \end{cases}$$



edge labels (shortest path values)

Unlabeled graphs

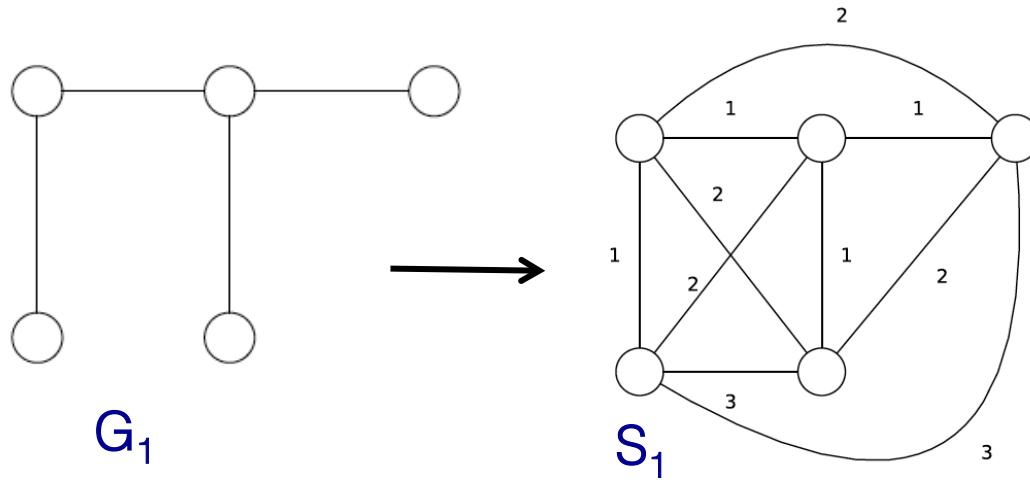
$$k_{\text{walk}}^{(1)} = \begin{cases} 1, & \text{if } \ell(e_1) = \ell(e_2) \text{ and } \ell(u_1) = \ell(u_2) \text{ and } \ell(v_1) = \ell(v_2) \\ 0, & \text{otherwise} \end{cases}$$

edge endpoint labels:
 $e_1 = (u_1, v_1)$

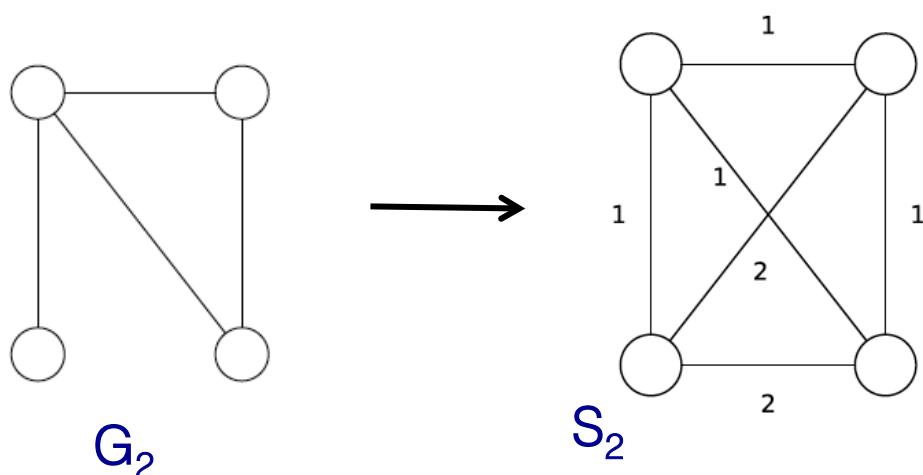
Labeled graphs



Shortest Path Kernels – Example



- In S_1 we have:
 - 4 edges with label 1
 - 4 edges with label 2
 - 2 edges with label 3
- In S_2 we have:
 - 4 edges with label 1
 - 2 edges with label 2



$$k(G_1, G_2) = \sum_{e_1 \in E_1} \sum_{e_2 \in E_2} k_{\text{walk}}^{(1)}(e_1, e_2)$$
$$= 4 \times 4 + 4 \times 2 = 24$$

Graph Similarity – What to Remember

- Numerous applications in graph mining and learning
 - Network monitoring, anomaly detection, network intrusion, behavioral studies
- Although seems easy problem, it's not
 - Some measures are counter-intuitive
 - Graph isomorphism is a computationally difficult problem
 - DeltaCon [Koutra et al. '13] (based on node proximity) satisfies several intuitive properties
- There are multiple measures, but which one to use?
 - Depends on the application!
 - Good news according to the guide of [Soundarajan et al. '14]
- Graph kernels and application in graph classification
 - Kernel trick: they can directly be applied in well known learning algorithms (e.g., SVMs)
 - Scalability issues

Thank You!

